

Comparing Mutual Information in Finite Samples: Wald and Expanded Welch–Satterthwaite Inference

Daniel Zhao

Supervisor: Associate Professor Joseph T. Lizier

A thesis submitted in partial fulfilment of the requirements for the degree of

Master of Computer Science

School of Computer Science

The University of Sydney

Australia

2026

Abstract

Mutual information (MI) measures association between two variables. This thesis asks whether two independently sampled populations have equal MI, even when their joint probability tables differ. The usual Normal Wald test uses a normal approximation that can be inaccurate in small or sparse samples.

The thesis develops an Expanded Welch–Satterthwaite adjustment for this comparison. It derives how changing one observation affects the estimated MI variance, including changes to the cell probabilities and their row and column totals. That calculation determines the degrees of freedom for a Student t distribution. Expanded Welch retains Wald’s estimated MI difference and standard error; it changes only the distribution used to judge the resulting statistic.

The methods are evaluated in 3,111 fixed statistical settings, each with 20,000 independently sampled table pairs. The design includes square tables from 2×2 to 8×8 , selected rectangular tables, unequal sample sizes, sparse cells, and populations with different distributions but equal MI. At a nominal 5% threshold, the experiment measures false positives, detection when MI differs, and the frequency of usable results.

Expanded Welch reduces excessive Wald false positives in some settings. However, when both tests return a p-value, its Student t distribution gives a p-value no smaller than Wald’s. It therefore also lowers rejection where Wald already rejects too rarely and can reduce detection when MI differs. In sparse samples, Expanded Welch more often fails to return a usable p-value under the implemented rule. Wald is closer to the target 5% rate in all 24 main null settings with sample sizes 500 or 1,000. A supplementary diagnostic shows that changing the reference distribution cannot correct error in the standard error shared by both tests. For the settings studied, Wald remains the simpler general baseline. Expanded Welch is a potential additional check when a separate, closely matched calibration study shows that Wald rejects equal MI too often.

Acknowledgements

I am grateful to my supervisor, Associate Professor Joseph T. Lizier, for guidance throughout this research. I also thank the colleagues, friends, and family members whose support made this work possible.

Contents

Abstract	i
Acknowledgements	ii
1 Introduction	1
1.1 Why the comparison is difficult	1
1.2 Research questions	2
1.3 Contributions to the field	2
1.4 Scope and organisation	3
2 Background and Literature Review	4
2.1 Information-theoretic foundations	4
2.1.1 Entropy and information content	4
2.1.2 Pointwise and average mutual information	5
2.2 Estimating discrete mutual information	6
2.3 Sampling variance of plug-in MI	7
2.4 Approaches to Testing MI Differences	11
2.5 Previous work and this thesis’s contribution	14
3 The Two-Sample MI Test and Wald Baseline	16
3.1 Problem setup and bias-corrected estimator	16
3.2 Shared standard error and test statistic	18
3.3 Normal Wald reference and motivation for expansion	18
4 The Expanded Welch–Satterthwaite Test for MI	20
4.1 Choosing Degrees of Freedom for Expanded Welch	20
4.2 How the MI Variance Changes with the Table	21

4.3	Component degrees of freedom and final test	25
4.4	Worked example and implementation	27
4.5	Behaviour and limitations	28
5	Experimental Design	30
5.1	A Concrete Pair of Population Tables	30
5.2	Constructing fixed joint distributions	30
5.3	Sampling, tests and measured outcomes	32
5.4	Experiment families and computation	33
6	Experimental Results and Discussion	35
6.1	Calibration across table size and margins	35
6.2	Detection as the true MI difference grows	39
6.3	Validity under sparse sampling	41
6.4	Checks on population and sampling choices	43
6.5	Limits of changing the reference distribution	43
6.6	Larger samples and computational cost	44
7	Implications and Limitations	47
7.1	Implications for using Expanded Welch	47
7.2	Limits of the evidence	48
8	Conclusion	49
A	Additional Derivation Details	51
A.1	Gradient of mutual information on the probability simplex	51
A.2	Why the variance of pointwise MI appears	52
A.3	Mean of the variance influence function	52
A.4	Direct table substitutions	53
A.5	Fixed-margin behaviour near independence	54
A.6	Dependence between the MI and variance estimates	55
A.7	Reference ordering and zero components	56
A.8	Numerical derivative checks	56
B	Population Construction and Exact Design Settings	57
B.1	Additive construction	57

B.2	Statistical experiment grids	58
C	Simulation Uncertainty and Validity Diagnostics	61
C.1	Rejection rates and simulation uncertainty	61
C.2	Validity and sparse-sample diagnostics	62
C.3	Independent-pilot standard-error check	62
D	Additional Experimental Figures	63
D.1	Remaining square-table comparisons	63
D.2	Other controlled comparisons	72
D.3	Main-grid null-rate overview	79
E	The Independence Boundary	80
E.1	Constructing the independent reference	80
E.2	Why the first-order term vanishes	81
E.3	Second-order behaviour	81
E.4	What alternative constructions produce	83

List of Figures

6.1	Rejection rates: 2×2 tables, uniform margins	36
6.2	Rejection rates: 8×8 tables, uniform margins	37
6.3	Rejection rates: 8×8 tables, different skew margins	38
6.4	Wider MI differences, uniform 2×2	40
6.5	Baseline-MI comparison, 3×3 tables	41
6.6	Extreme-skew 3×3 tables	42
6.7	Large-sample null behaviour	45
D.1	Rejection rates: 2×2 tables, same skew margins	64
D.2	Rejection rates: 2×2 tables, different skew margins	65
D.3	Rejection rates: 3×3 tables, uniform margins	66
D.4	Rejection rates: 3×3 tables, same skew margins	67
D.5	Rejection rates: 3×3 tables, different skew margins	68
D.6	Rejection rates: 5×5 tables, uniform margins	69
D.7	Rejection rates: 5×5 tables, same skew margins	70
D.8	Rejection rates: 5×5 tables, different skew margins	71
D.9	Rejection rates: 8×8 tables, same skew margins	72
D.10	Changed-cell comparison, 5×5 tables	73
D.11	Unequal samples, Q higher	74
D.12	Unequal samples, P higher	75
D.13	Rectangular tables	76
D.14	Additive versus log-linear construction	77
D.15	Exact-independence boundary	78

List of Tables

4.1	Operational Expanded Welch–Satterthwaite algorithm.	28
5.1	Row and column margins used for each population in the principal design. “Uniform” means the vector with every entry equal to $1/k$	32
5.2	The fixed experiment families. All statistical points use 20,000 paired replicates and the 5% threshold.	34
6.1	Rejection rates and paired differences in two fully valid null regimes.	39
6.2	Median method-call times and interquartile ranges, in microseconds.	46
B.1	Exact statistical experiment grids; MI values are in nats.	58
D.1	Main-grid null configurations below, inside or above the descriptive rejection-rate interval 0.025–0.075.	79

Chapter 1

Introduction

Suppose two independent groups are observed on the same categorical variables X and Y . Mutual information (MI) describes the strength of association between those variables in each group. If the estimated MI differs, the difference alone does not tell us whether the population MI values differ: both estimates vary when new observations are drawn. The question in this thesis is whether the two *population* MI values can plausibly be equal.

This is not the usual test of whether X and Y are independent within one group. The null hypothesis here permits dependence in both groups and permits their complete joint distributions to differ. It specifies only

$$H_0 : I(P) = I(Q), \quad H_1 : I(P) \neq I(Q), \quad (1.1)$$

where P and Q are the two joint distributions. The samples from P and Q are independent of each other; that sampling assumption does not imply independence of X and Y within either population.

1.1 Why the comparison is difficult

MI uses the probabilities in every cell of a table and the totals of its rows and columns. In a small sample, some cells can be empty even when their population probabilities are positive. The observed tables determine both the estimated MI difference and its standard error, so sparse data can affect both. The Normal Wald test compares their ratio with a standard normal distribution, but in small samples it may reject equal MI more or less often than the intended 5%.

Expanded Welch offers a possible adjustment. It uses the same estimated difference and standard error as Wald, but compares their ratio with a Student t

distribution. Its degrees of freedom depend on how much the estimated MI variance changes from sample to sample. To calculate that change, we must account for the fact that changing one cell also changes its row and column totals and the MI values calculated from them. This thesis derives that calculation and tests whether the adjustment helps.

The adjustment is not automatically better. It can reduce false positives where Wald rejects too often, but it can also reduce detection where Wald already rejects too rarely. In sparse samples, it may fail to return a usable p-value when Wald succeeds. The experiment therefore examines false positives, detection and the frequency of usable results together.

1.2 Research questions

The thesis asks two main questions. First, how can changes in the estimated MI variance be used to choose the degrees of freedom for a test using a Student t distribution? Second, does the resulting Expanded Welch test offer a useful improvement over Normal Wald? Answering the second question requires looking at false positives, detection when MI differs, and the frequency of usable results. The experiment checks how these outcomes change with table size, uneven cell probabilities, the amount and location of association, unequal sample sizes and larger samples. It also measures the extra computation required.

1.3 Contributions to the field

This thesis combines established ideas about MI and Welch–Satterthwaite tests for a specific question: whether two populations have equal MI. Its contributions are the calculation needed for that test and evidence about whether the test is useful.

1. **Mathematical contribution.** The thesis derives how changing one observation affects the estimated MI variance. The calculation includes the resulting changes in the row and column totals and in pointwise MI.
2. **Methodological contribution.** Expanded Welch uses this calculation to choose the degrees of freedom for a two-sample test using a Student t distribution. The main experiment compares it with Wald; a supplementary analysis checks simpler sample-size choices for the degrees of freedom.
3. **Empirical contribution.** A reproducible experiment measures false positives, detection, usable results and runtime on fixed populations with known MI. It

covers 2×2 through 8×8 square tables, rectangular tables, unequal samples and very sparse settings.

4. **Practical contribution.** Expanded Welch can help when Wald rejects equal MI too often. It can also make an already-low rejection rate lower and return fewer usable results in sparse samples. The evidence shows why it should not replace Wald in general and points to problems that changing the reference distribution alone cannot solve.

1.4 Scope and organisation

The analysis concerns independent samples recorded in count tables with the same row and column categories. The main derivation assumes every population cell has positive probability and that its leading large-sample MI variance is greater than zero. Exact independence is considered separately because that variance becomes zero. The experiment compares Normal Wald with Expanded Welch using a two-sided 5% test. It does not establish performance for paired data, continuous MI estimators or a number of categories that grows with sample size.

Chapter 2 supplies the statistical background. Chapters 3 and 4 define the shared statistic and derive the correction. Chapter 5 explains how the populations and samples are generated and how the results are measured. Chapter 6 presents and interprets the main comparisons in sequence. Chapter 7 draws out their broader implications and limits, and Chapter 8 states the contribution and recommendation. The appendices give the detailed algebra, population settings, simulation diagnostics and supporting figures.

Chapter 2

Background and Literature Review

This chapter develops the concepts needed to compare MI between two categorical populations. It begins with entropy, pointwise MI, and population MI. It then reviews plug-in estimation, first-order inference, the Welch–Satterthwaite approximation, and prior comparisons of information quantities. The final section identifies the gap addressed by the expanded method.

2.1 Information-theoretic foundations

2.1.1 Entropy and information content

For a discrete random variable X with probability mass function $p_X(x)$, the information content of outcome x is $-\log p_X(x)$. Rare outcomes have larger information content because they are less expected. Using natural logarithms gives information in nats. The Shannon entropy is the expected information content,

$$H(X) = - \sum_x p_X(x) \log p_X(x). \quad (2.1)$$

Entropy summarises the uncertainty of one random variable before its value is observed ([Shannon, 1948](#)).

For a pair (X, Y) , the joint entropy $H(X, Y)$ is defined from the joint probabilities in the same way. The conditional entropy $H(X | Y) = H(X, Y) - H(Y)$ measures the uncertainty remaining in X after observing Y .

2.1.2 Pointwise and average mutual information

Let $P = (p_{ij})$ be the joint distribution of categorical variables $X \in \{1, \dots, r\}$ and $Y \in \{1, \dots, c\}$. Write

$$p_{i+} = \sum_{j=1}^c p_{ij}, \quad p_{+j} = \sum_{i=1}^r p_{ij} \quad (2.2)$$

for the row and column margins. The pointwise mutual information associated with cell (i, j) is

$$\ell_P(i, j) = \log \left(\frac{p_{ij}}{p_{i+}p_{+j}} \right). \quad (2.3)$$

The denominator is the cell probability under independence with the same margins. A positive value means that the cell occurs more often than independence predicts; a negative value means that it occurs less often.

Mutual information is the probability-weighted average of these pointwise values (Shannon, 1948; Bossomaier et al., 2016):

$$I(P) = \sum_{i=1}^r \sum_{j=1}^c p_{ij} \ell_P(i, j). \quad (2.4)$$

Here the summations are nested: for each row i , the inner sum runs over all columns j . Equivalently, the expression sums once over every cell in the table and may be written with the shorthand $\sum_{i,j}$. Because P is the joint distribution of (X, Y) , the same quantity is conventionally written as $I(X; Y)$. Equivalently,

$$I(X; Y) = H(X) + H(Y) - H(X, Y). \quad (2.5)$$

The quantity is nonnegative and equals zero exactly when $p_{ij} = p_{i+}p_{+j}$ for every cell, including zero-probability cells (Bossomaier et al., 2016).

MI measures the overall departure from independence, not one particular shape of association.

This scalar summary deliberately discards much of the information in the joint distribution. Two populations can have the same MI while assigning probability to different cells, having different margins, or exhibiting different association patterns. Equality of MI is therefore weaker than equality of distributions. Our null requires only this one number to agree, not that the two tables match cell by cell.

The logarithm base determines the reporting unit. This thesis uses natural logarithms, so MI is measured in nats; dividing by $\log 2$ converts to bits. This conversion rescales MI and its standard error together but leaves the degrees of freedom unchanged, so the test result is identical in either unit.

2.2 Estimating discrete mutual information

Suppose a sample of size n produces counts N_{ij} . The empirical cell probabilities are

$$\hat{p}_{ij} = \frac{N_{ij}}{n}. \quad (2.6)$$

Substitution into Equation (2.4) gives the plug-in estimator

$$\hat{I}(P) = \sum_{i,j:\hat{p}_{ij}>0} \hat{p}_{ij} \log \left(\frac{\hat{p}_{ij}}{\hat{p}_{i+}\hat{p}_{+j}} \right). \quad (2.7)$$

The zero-cell convention follows from $\lim_{x \downarrow 0} x \log x = 0$. Empty cells contribute no direct term, but they still affect the empirical margins and the finite-sample behaviour of the estimator.

Under multinomial sampling, the cell counts are random but constrained to sum to n , so they are dependent. The plug-in estimator remains consistent for a fixed alphabet, but its finite-sample behaviour depends on the complete probability table. In particular, rare cells can change the sampling behaviour even when the alphabet and sample sizes are unchanged (Paninski, 2003).

For a fixed positive $r \times c$ support, the plug-in estimator has a positive leading bias as the sample size grows:

$$\text{Bias}\{\hat{I}(P)\} = \frac{(r-1)(c-1)}{2n} + O(n^{-2}). \quad (2.8)$$

The leading term grows with table size and shrinks as the sample grows. It does not describe every small or sparse sample. The result is closely related to Miller's entropy correction (Miller, 1955), and the same histogram-MI term appears in Moddemeijer (1989). Sparse sampling can produce further error beyond this leading approximation (Paninski, 2003).

Bias matters directly when comparing two populations. If $I(P) = I(Q)$ but $n_P \neq n_Q$, the uncorrected difference has leading bias

$$\frac{(r-1)(c-1)}{2} \left(\frac{1}{n_P} - \frac{1}{n_Q} \right). \quad (2.9)$$

For a 2×2 table with $n_P = 50$ and $n_Q = 100$, Equation (2.9) gives 0.005 nats despite equal population MIs. Chapter 3 corrects this sample-size effect by subtracting $(r-1)(c-1)/(2n)$ from each estimate.

2.3 Sampling variance of plug-in MI

To compare two estimated MI values, it is necessary to know how much each estimate would vary across repeated samples. The goal of this section is to derive the large-sample approximation

$$\text{Var}\{\hat{I}(P)\} \approx \frac{V(P)}{n}. \quad (2.10)$$

Here $V(P)$, defined below, measures how much the contribution of one observation varies across the population. The derivation assumes that the population cell probabilities are positive, so the logarithms and their derivatives are well defined. This first-order variance is established for histogram and discrete MI (Moddemeijer, 1989; Brillinger, 2004). The argument proceeds in three steps. First, sampling changes the empirical cell probabilities. Second, a first-order Taylor expansion translates those probability errors into an MI error. Third, that error can be written as a sample average, whose variance is straightforward to calculate.

Step 1: From probability errors to MI error.

Using the entropy identity in Equation (2.5), MI can be written directly in terms of the joint and marginal probabilities:

$$I(P) = \sum_{i,j} p_{ij} \log p_{ij} - \sum_i p_{i+} \log p_{i+} - \sum_j p_{+j} \log p_{+j}. \quad (2.11)$$

Let $f(u) = u \log u$. The derivative needed for its first-order Taylor expansion follows directly from the product rule:

$$f'(u) = \frac{d}{du} \{u \log u\} = \log u + u \frac{1}{u} = \log u + 1. \quad (2.12)$$

The general first-order Taylor formula is

$$f(x) \approx f(a) + f'(a)(x - a), \quad (2.13)$$

where a is the value around which the function is expanded and $x - a$ is the distance from that value. Here the new value is the estimated probability $x = \hat{u}$, the expansion point is the population probability $a = u$, and $f(t) = t \log t$. Therefore, $x - a = \hat{u} - u$, which is the sampling error in the estimated probability. Substituting these quantities into the general formula gives

$$\hat{u} \log \hat{u} \approx \underbrace{u \log u}_{\text{population value}} + \underbrace{(\log u + 1)(\hat{u} - u)}_{\text{first-order change}}. \quad (2.14)$$

Subtracting the population value from both sides makes the change explicit:

$$\hat{u} \log \hat{u} - u \log u \approx (\log u + 1)(\hat{u} - u). \quad (2.15)$$

Applying this expansion to the three sums in Equation (2.11) gives

$$\begin{aligned} \hat{I}(P) - I(P) &\approx \sum_{i,j} (\log p_{ij} + 1)(\hat{p}_{ij} - p_{ij}) \\ &\quad - \sum_i (\log p_{i+} + 1)(\hat{p}_{i+} - p_{i+}) \\ &\quad - \sum_j (\log p_{+j} + 1)(\hat{p}_{+j} - p_{+j}). \end{aligned} \quad (2.16)$$

The margin errors are sums of the corresponding cell errors:

$$\hat{p}_{i+} - p_{i+} = \sum_j (\hat{p}_{ij} - p_{ij}), \quad \hat{p}_{+j} - p_{+j} = \sum_i (\hat{p}_{ij} - p_{ij}). \quad (2.17)$$

For example, the row-margin sum can now be written as

$$\begin{aligned} \sum_i (\log p_{i+} + 1)(\hat{p}_{i+} - p_{i+}) &= \sum_i \left[(\log p_{i+} + 1) \left\{ \sum_j (\hat{p}_{ij} - p_{ij}) \right\} \right] \\ &= \sum_{i,j} (\log p_{i+} + 1)(\hat{p}_{ij} - p_{ij}). \end{aligned} \quad (2.18)$$

The inner sum over j is the probability error for a fixed row i . It is multiplied by that row's coefficient $\log p_{i+} + 1$; the outer sum then adds these products over all rows. The column-margin sum can be rearranged in the same way. Consequently, the same cell error $\hat{p}_{ij} - p_{ij}$ appears in the joint term, its row-margin term and its column-margin term. Its combined coefficient is

$$\begin{aligned} &(\log p_{ij} + 1) - (\log p_{i+} + 1) - (\log p_{+j} + 1) \\ &= \log \left(\frac{p_{ij}}{p_{i+} p_{+j}} \right) - 1 \\ &= \ell_P(i, j) - 1. \end{aligned} \quad (2.19)$$

Therefore,

$$\hat{I}(P) - I(P) \approx \sum_{i,j} \{\ell_P(i, j) - 1\}(\hat{p}_{ij} - p_{ij}). \quad (2.20)$$

The empirical and population cell probabilities both sum to one, so

$$\sum_{i,j} (\hat{p}_{ij} - p_{ij}) = \sum_{i,j} \hat{p}_{ij} - \sum_{i,j} p_{ij} = 1 - 1 = 0. \quad (2.21)$$

The constant -1 in Equation (2.19) therefore contributes zero. Explicitly,

$$\begin{aligned}
 \hat{I}(P) - I(P) &\approx \sum_{i,j} \{\ell_P(i,j) - 1\}(\hat{p}_{ij} - p_{ij}) \\
 &= \sum_{i,j} \ell_P(i,j)(\hat{p}_{ij} - p_{ij}) - \sum_{i,j} (\hat{p}_{ij} - p_{ij}) \\
 &= \sum_{i,j} \ell_P(i,j)(\hat{p}_{ij} - p_{ij}). \tag{2.22}
 \end{aligned}$$

The complicated changes in the joint table and its two margins have therefore reduced to a weighted sum of cell-probability errors. The weight for a cell is its population pointwise MI.

Step 2: From cell errors to observation-level errors.

Let $Z_a = (X_a, Y_a)$ be observation a . An empirical cell probability is the proportion of observations in that cell, so

$$\hat{p}_{ij} - p_{ij} = \frac{1}{n} \sum_{a=1}^n \{\mathbf{1}(Z_a = (i,j)) - p_{ij}\}. \tag{2.23}$$

Substituting this expression into Equation (2.22) first gives

$$\hat{I}(P) - I(P) \approx \sum_{i,j} \ell_P(i,j) \left[\frac{1}{n} \sum_{a=1}^n \{\mathbf{1}(Z_a = (i,j)) - p_{ij}\} \right]. \tag{2.24}$$

The finite sums can be exchanged, so the contribution from observation a is gathered inside one pair of brackets:

$$\begin{aligned}
 \hat{I}(P) - I(P) &\approx \frac{1}{n} \sum_{a=1}^n \left[\sum_{i,j} \ell_P(i,j) \mathbf{1}(Z_a = (i,j)) - \sum_{i,j} p_{ij} \ell_P(i,j) \right] \\
 &= \frac{1}{n} \sum_{a=1}^n \{\ell_P(Z_a) - I(P)\}. \tag{2.25}
 \end{aligned}$$

The first sum inside the brackets selects the single cell containing Z_a , giving $\ell_P(Z_a)$. The second is the probability-weighted average of all pointwise-MI values, which is $I(P)$. This equation is the key step. To first order, the error in the plug-in MI estimate behaves like the average of n observation-level contributions. The contribution of an observation in cell (i,j) is simply that cell's pointwise MI minus the population average, $\ell_P(i,j) - I(P)$. Its mean is zero because

$$E_P\{\ell_P(X,Y) - I(P)\} = \sum_{i,j} p_{ij} \{\ell_P(i,j) - I(P)\}$$

$$= \underbrace{\sum_{i,j} p_{ij} \ell_P(i, j)}_{I(P)} - I(P) \underbrace{\sum_{i,j} p_{ij}}_1 = 0. \quad (2.26)$$

In formal asymptotic terminology, the centred function $\ell_P(i, j) - I(P)$ is the influence function of MI; Equation (2.25) is its first-order linear representation (Hampel, 1974; van der Vaart, 1998; Kandasamy et al., 2015).

Step 3: Calculate the sampling variance.

Define the population variance of these centred contributions by

$$V(P) = \text{Var}_P\{\ell_P(X, Y) - I(P)\} = \sum_{i,j} p_{ij} \{\ell_P(i, j) - I(P)\}^2. \quad (2.27)$$

Under independent sampling, the n terms in Equation (2.25) are independent and each has variance $V(P)$. The covariance between any two different observations is therefore zero. Hence

$$\begin{aligned} \text{Var}\{\hat{I}(P)\} &\approx \text{Var}\left\{\frac{1}{n} \sum_{a=1}^n \{\ell_P(Z_a) - I(P)\}\right\} \\ &= \frac{1}{n^2} \sum_{a=1}^n \text{Var}\{\ell_P(Z_a) - I(P)\} \\ &= \frac{1}{n^2} \sum_{a=1}^n V(P) = \frac{nV(P)}{n^2} = \frac{V(P)}{n}. \end{aligned} \quad (2.28)$$

This proves the approximation stated in Equation (2.10). It is the familiar result that the variance of an average of n independent contributions is the variance of one contribution divided by n .

The independence of observations is needed in the second line of Equation (2.28): it makes covariance terms between different observations vanish. Moddemeijer (1999) studies MI sampling variance when observation pairs are dependent and those terms must instead be accounted for. The present test and experiment concern independent multinomial samples, so they use $V(P)/n$ without a serial-dependence correction.

When $V(P) > 0$, the central limit theorem also gives

$$\frac{\hat{I}(P) - I(P)}{\sqrt{V(P)/n}} \xrightarrow{d} N(0, 1). \quad (2.29)$$

This does not claim that every finite-sample MI estimate is normally distributed. It is a local, large-sample approximation for a fixed table with positive population cell probabilities. The $1/n$ leading bias from Section 2.2 decreases faster than the $1/\sqrt{n}$ sampling error, so it does not change this limiting distribution.

The independence boundary.

At exact independence, $p_{ij} = p_{i+}p_{+j}$, so $\ell_P(i, j) = \log 1 = 0$ in every cell. Hence $I(P) = 0$, $V(P) = 0$, and the first-order average in Equation (2.25) vanishes.

A finite sample can still have positive MI. Its leading variation comes from the squared cell-probability errors in a second-order Taylor expansion, which produces the usual chi-squared limiting distribution (Brillinger, 2004; Marinescu and Balcau, 2025). The normal limiting distribution for comparing positive MI values and the chi-squared limiting distribution for testing independence therefore apply in different settings.

2.4 Approaches to Testing MI Differences

The methods below answer related but different questions. Some quantify the sampling error of one MI estimate; others compare two estimates or adjust a test when its standard error is itself estimated. For this thesis, a useful predecessor must be assessed against the null $I(P) = I(Q)$, which allows the two population tables to differ.

Normal Wald inference.

A Wald test estimates a parameter difference, divides it by an estimated standard error, and compares the result with a standard normal distribution (Wald, 1943; van der Vaart, 1998). For two independent MI estimates, the two first-order variance contributions add:

$$\text{Var}\{\hat{I}(P) - \hat{I}(Q)\} \approx \frac{V(P)}{n_P} + \frac{V(Q)}{n_Q}. \quad (2.30)$$

Each term is the sampling variance of one MI estimate. Independence of the two samples makes their covariance zero. Replacing the population variances by estimates gives the standardisation

$$\frac{\hat{I}(P) - \hat{I}(Q)}{\sqrt{\hat{V}(P)/n_P + \hat{V}(Q)/n_Q}} \xrightarrow{d} N(0, 1) \quad \text{under } I(P) = I(Q). \quad (2.31)$$

Here $\hat{V}(P)$ and $\hat{V}(Q)$ estimate the pointwise-MI variances $V(P)$ and $V(Q)$. This limit assumes the regular positive-support conditions of Section 2.3 and increasing sample sizes. Chapter 3 applies the leading bias correction to the numerator and specifies the full Normal Wald test.

With smaller samples, the MI difference may not follow the normal approximation closely, and the estimated standard error can vary between samples. The Wald test

still uses a normal distribution; it makes no separate small-sample adjustment for that standard-error variability.

The Welch–Satterthwaite principle.

Welch’s test compares two means when their population variances may differ (Welch, 1947). For two independent samples with means $\bar{X}_P, \bar{X}_Q$ and variances s_P^2, s_Q^2 , the statistic for testing equal population means and its effective degrees of freedom are

$$T_{\text{Welch}} = \frac{\bar{X}_P - \bar{X}_Q}{\sqrt{s_P^2/n_P + s_Q^2/n_Q}}, \quad (2.32)$$

$$\nu = \frac{\left(s_P^2/n_P + s_Q^2/n_Q\right)^2}{\frac{(s_P^2/n_P)^2}{n_P - 1} + \frac{(s_Q^2/n_Q)^2}{n_Q - 1}}. \quad (2.33)$$

The denominator of the statistic contains one estimated variance contribution from each sample. Under normal sampling, their component degrees of freedom are $n_P - 1$ and $n_Q - 1$. If s_P^2/n_P dominates, the formula makes the combined ν close to $n_P - 1$; when both contributions matter, both component values affect ν . Satterthwaite obtains this form by matching the mean and variance of their sum to a scaled chi-squared variable (Satterthwaite, 1946).

Welch compares the statistic with a Student t distribution using the calculated ν . At the same significance level, this requires a larger absolute statistic to reject the null than a standard normal distribution does. The difference between the rejection thresholds from the Student t and standard normal distributions becomes smaller as ν increases.

For MI, the estimated variance contributions are $\hat{V}(P)/n_P$ and $\hat{V}(Q)/n_Q$. Their component degrees of freedom need their own calculation because $\hat{V}(P)$ and $\hat{V}(Q)$ are nonlinear functions of the probability tables. Chapter 4 derives these MI-specific inputs to the same moment-matching structure.

Regression studies have also adjusted reference distributions to account for uncertainty in estimated variances (Pan and Wall, 2002; Bell and McCaffrey, 2002; Kauermann and Carroll, 2001; Fay and Graubard, 2001; Imbens and Kolesár, 2016). For example, Pan and Wall (2002) calculate approximate degrees of freedom by examining how a regression variance estimate varies between samples. This supports examining the variance estimate itself, but the estimated MI variance is calculated differently and needs its own analysis. Bell and McCaffrey (2002) report regression variance estimates biased downward, while Kauermann and Carroll (2001) examine intervals that cover too rarely. A more cautious reference can help in such settings.

In the MI comparison examined in Section 6.5, the average squared plug-in standard error exceeds the measured finite-sample variance of the estimated difference. Using a Student t distribution does not correct that discrepancy and lowers rejection further in that setting.

Prior MI inference and related comparisons.

MI uncertainty also has a broader literature than its leading normal approximation. [Roulston \(1999\)](#) derives entropy and MI error estimates using error propagation and examines them numerically. [Hutter \(2002\)](#) and [Hutter and Zaffalon \(2005\)](#) use a Bayesian model for the unknown cell probabilities to calculate uncertainty about MI after observing a table, including its variance and distributional shape. That differs from the repeated-sample question here: how often a test comparing two MI values rejects. Their results therefore do not directly justify using a Student t distribution for the test in this thesis.

Hutcheson’s test is a direct information-theoretic Welch predecessor. It compares two Shannon diversity estimates with an estimated variance and a Student t distribution, using

$$\nu_H = \frac{(s_{H,P}^2 + s_{H,Q}^2)^2}{(s_{H,P}^2)^2/n_P + (s_{H,Q}^2)^2/n_Q} \quad (2.34)$$

as its effective degrees of freedom ([Hutcheson, 1970](#)). Here each of $s_{H,P}^2$ and $s_{H,Q}^2$ is the estimated sampling variance of the corresponding entropy estimator $\widehat{H}$, and each variance component is assigned its sample size, not $n_i - 1$. Applying the same assignment to the two MI variance contributions gives a clear historical baseline. It does not, however, derive how the plug-in MI variance itself fluctuates: MI depends on a joint distribution and both margins, so moving mass into one cell changes pointwise values elsewhere in its row and column. The expanded calculation retains these changes.

The question of comparing MI values also has predecessors. In a segregation application, [Mora and Ruiz-Castillo \(2009\)](#) develop inference for an MI index and discuss pairwise comparisons across groups and periods; the later peer-reviewed account places that MI index within the wider segregation literature ([Mora and Ruiz-Castillo, 2011](#)). They do not derive the variance-sensitivity degrees of freedom used here. [Guha and Chothia \(2014\)](#) use MI for a different two-sample question: whether two continuous distributions are identical. Their MI is between a pooled observation and its sample label; their null is equality of distributions, not equality of two within-population MI values. For discrete data, [Drake and Guha \(2014\)](#) develop an MI-based k -sample test for the same distributional-equality question.

At exact independence, the first-order sampling variation of MI vanishes and

its second-order variation becomes important (Marinescu and Balcau, 2025). The positive-MI comparison developed here is therefore not a replacement for a one-table independence test. Second-order effects can also matter near that boundary.

Resampling alternatives.

A bootstrap would draw many new samples from the observed tables to estimate how much the test statistic varies. This requires repeating the calculation many times (Efron, 1979). A simple permutation test instead combines the two samples and shuffles their group labels. That shuffle is justified when the populations have the same complete distribution, $P = Q$. Our hypothesis requires only $I(P) = I(Q)$: the tables can differ while their MI values are equal. Shuffling labels could therefore test a stronger claim than this thesis asks about (Chung and Romano, 2013). More specialised permutation methods are possible for some such comparisons (Janssen, 1997; Chung and Romano, 2013), but assessing one for equal MI would be a separate study. The present comparison focuses on Normal Wald and Expanded Welch.

2.5 Previous work and this thesis’s contribution

The cited work supplies several parts of the present comparison. Moddemeyer (1989) and Brillinger (2004) give the leading sampling variance of an MI estimate; Miller (1955) and Paninski (2003) explain why finite-sample bias and sparse tables also matter. Hutcheson’s entropy comparison uses a Student t distribution with degrees of freedom assigned from sample sizes (Hutcheson, 1970). In a different setting, Pan and Wall (2002) show how uncertainty in a non-classical variance estimate can inform degrees of freedom. These are useful precedents, but none supplies the change in the plug-in MI variance when its observed table changes.

The question of comparing information values also has precedents. Mora and Ruiz-Castillo (2009) discusses comparisons of an MI-based segregation index. Tests based on MI between pooled observations and their group labels instead target equality of the complete distributions (Guha and Chothia, 2014; Drake and Guha, 2014). The present null is narrower: $I(P) = I(Q)$ permits $P \neq Q$. This difference in the null hypothesis determines which existing reference distributions can be used directly.

The remaining calculation is to measure how much the *estimated MI variance* changes from sample to sample. Assigning each population degrees of freedom from its sample size does not measure this change. A new observation affects a cell probability, its row and column totals, and the MI values calculated from them. This thesis derives how those changes affect the estimated variance and uses the

result to choose a Student t distribution for comparing two MI values. It then tests the method on fixed population tables, showing when it brings false-positive rates closer to the target and when it instead rejects too rarely.

Chapter 3 defines the comparison, Chapter 4 derives the adjustment, and Chapter 6 evaluates it.

Chapter 3

The Two-Sample MI Test and Wald Baseline

This chapter defines the population MI difference being tested and the statistic calculated from the observed tables. Both methods use this same statistic. Normal Wald judges it against a standard normal distribution; Chapter 4 develops a Student t distribution with estimated degrees of freedom.

3.1 Problem setup and bias-corrected estimator

Let

$$N^P = (N_{ij}^P), \quad N^Q = (N_{ij}^Q) \quad (3.1)$$

be independent $r \times c$ count tables with totals n_P and n_Q . The row and column categories have the same definitions in both populations. The empirical probabilities are

$$\hat{p}_{ij} = \frac{N_{ij}^P}{n_P}, \quad \hat{q}_{ij} = \frac{N_{ij}^Q}{n_Q}. \quad (3.2)$$

The population parameter is the signed difference

$$\Delta = I(P) - I(Q). \quad (3.3)$$

The principal test is two-sided:

$$H_0 : \Delta = 0 \quad \text{against} \quad H_1 : \Delta \neq 0. \quad (3.4)$$

The large-sample argument uses the following assumptions.

Assumption 1 (Sampling). *Observations are independent and identically dis-*

tributed within each population, and the two population samples are independent.

Assumption 2 (Aligned fixed alphabets). *The row and column categories are fixed in advance and aligned between the two tables.*

Assumption 3 (Positive cell probabilities). *Every cell in the population table has a probability greater than zero. A sample can still contain no observations in a cell, especially when that cell is rare. Such samples remain in the study; we record how often cells have zero observed counts to show how sparse the samples are.*

Assumption 4 (Positive MI). *Both populations have positive MI: $I(P) > 0$ and $I(Q) > 0$. With the positive cell probabilities in Assumption 3, this is equivalent to requiring positive first-order variances $V(P) > 0$ and $V(Q) > 0$. It excludes exact independence, where the first-order approximation breaks down.*

Assumption 5 (Logarithm units). *All logarithms in this thesis are natural, so MI differences and standard errors are reported in nats. Changing the logarithm base scales both by the same positive constant and leaves the standardised statistic unchanged. This is a choice of units, not a condition for the test to be valid.*

Bias-corrected estimator.

For the first table, define empirical pointwise MI by

$$\hat{\ell}_P(i, j) = \log \left(\frac{\hat{p}_{ij}}{\hat{p}_{i+}\hat{p}_{+j}} \right) \quad \text{when } \hat{p}_{ij} > 0. \quad (3.5)$$

The plug-in MI estimate is

$$\hat{I}(P) = \sum_{i,j:\hat{p}_{ij}>0} \hat{p}_{ij} \hat{\ell}_P(i, j). \quad (3.6)$$

The same calculation gives $\hat{I}(Q)$.

The leading fixed-support MI bias appears in Equation (2.8) (Moddemeijer, 1989). Subtracting it from each plug-in estimate gives the estimated effect:

$$\hat{\Delta}_{\text{BC}} = \left\{ \hat{I}(P) - \frac{(r-1)(c-1)}{2n_P} \right\} - \left\{ \hat{I}(Q) - \frac{(r-1)(c-1)}{2n_Q} \right\}. \quad (3.7)$$

The two terms in braces are the bias-corrected MI estimates for the separate populations.

The correction uses the declared alphabet dimensions, not the observed nonempty dimensions. A corrected estimate can be negative in a sparse sample and is not clipped at zero, since clipping would define a different nonlinear estimator. When

$n_P = n_Q$, the two corrections cancel in the difference; otherwise, their difference removes the leading unequal-sample bias described in Section 2.2.

3.2 Shared standard error and test statistic

Section 2.3 showed that the first-order sampling variance of $\hat{I}(P)$ is $V(P)/n_P$, where $V(P)$ is the population variance of pointwise MI (Moddemeijer, 1989; Brillinger, 2004). Subtracting the fixed leading bias term in Equation (3.7) does not change this variance.

The observed table estimates $V(P)$ by

$$\hat{V}(P) = \sum_{i,j:\hat{p}_{ij}>0} \hat{p}_{ij} \{\hat{\ell}_P(i,j) - \hat{I}(P)\}^2. \quad (3.8)$$

Define $\hat{V}(Q)$ analogously.

Because the two samples are independent, their estimated variance contributions add:

$$\widehat{\text{SE}}^2 = \frac{\hat{V}(P)}{n_P} + \frac{\hat{V}(Q)}{n_Q}. \quad (3.9)$$

The common standardised statistic is

$$T = \frac{\hat{\Delta}_{\text{BC}}}{\widehat{\text{SE}}}. \quad (3.10)$$

The null hypothesis is not imposed when calculating $\hat{\Delta}_{\text{BC}}$, $\widehat{\text{SE}}$ or T : they use the observed tables. Equal population MI enters when interpreting the reference distribution, which is centred at zero under H_0 . Equations (3.7)–(3.10) are shared by Normal Wald and Expanded Welch. Their difference lies in the reference distribution, not the effect estimate or plug-in standard error.

3.3 Normal Wald reference and motivation for expansion

When both sample sizes grow while the table dimensions and positive population probabilities stay fixed, the statistic approaches a standard normal distribution under equal population MI (Wald, 1943; van der Vaart, 1998):

$$T \xrightarrow{d} N(0, 1) \quad \text{under } H_0. \quad (3.11)$$

The two-sided Normal Wald p-value is

$$p_{\text{Wald}} = 2\{1 - \Phi(|T|)\}, \quad (3.12)$$

where Φ is the standard normal distribution function.

Using the standard normal distribution is a large-sample approximation here. Wald estimates the standard error from the tables but applies no degrees-of-freedom adjustment.

Why consider an expanded reference?

The shared statistic supplies the estimated effect and its first-order standard error. Earlier information-theoretic and ordinary-Welch analogues set each variance component's degrees of freedom from its sample size: n_i (Hutcheson, 1970) or $n_i - 1$ (Welch, 1947). Neither assignment measures how much the plug-in MI variance changes when the observed table changes. That variance depends on joint cells, margins and pointwise MI values. Expanded Welch instead estimates each component's degrees of freedom from how the complete MI variance estimate changes with the table.

To choose the degrees of freedom, we need the average value of each estimated MI variance and how much that estimate varies across samples (Satterthwaite, 1946). For P , to first order,

$$E\{\hat{V}(P)\} \approx V(P). \quad (3.13)$$

Chapter 4 derives $\text{Var}\{\hat{V}(P)\}$ by tracking how the calculated MI variance changes when one observation changes the joint table.

Chapter 4

The Expanded Welch–Satterthwaite Test for MI

Expanded Welch keeps the estimated MI difference, standard error and test statistic from Chapter 3. It changes the distribution used to judge that statistic from the standard normal distribution to a Student t distribution. To choose its degrees of freedom, this chapter calculates how much the estimated MI variance changes between samples.

4.1 Choosing Degrees of Freedom for Expanded Welch

For the two MI estimates, the Welch–Satterthwaite equation combines their estimated variance contributions as in the original moment-matching framework (Satterthwaite, 1946; Welch, 1947):

$$\nu = \frac{\left\{ \hat{V}(P)/n_P + \hat{V}(Q)/n_Q \right\}^2}{\left\{ \hat{V}(P)/n_P \right\}^2 / \nu_P + \left\{ \hat{V}(Q)/n_Q \right\}^2 / \nu_Q}. \quad (4.1)$$

The numerator contains the two terms in the squared standard error. The remaining quantities are the component degrees of freedom ν_P and ν_Q , which should measure how reliably $\hat{V}(P)$ and $\hat{V}(Q)$ estimate their population counterparts.

The derivation has four steps. First, we relate the degrees of freedom to the average and across-sample variation of $\hat{V}(P)$. Second, we calculate how $V(P)$ changes when a small amount of probability moves into one cell. Third, that change tells us how variable $\hat{V}(P)$ is. Finally, we combine the degrees of freedom from P and Q using Equation (4.1).

Moment matching for one component.

Satterthwaite represents a positive variance estimate by a scaled chi-squared variable. For the first population, the working approximation represents $\widehat{V}(P)$ by

$$\mathbb{E}\{\widehat{V}(P)\} \frac{\chi_{\nu_P}^2}{\nu_P}. \quad (4.2)$$

The approximation matches moments; it does not assert that the finite-sample distribution of $\widehat{V}(P)$ is exactly chi-squared.

For a chi-squared variable with ν_P degrees of freedom,

$$\mathbb{E}(\chi_{\nu_P}^2) = \nu_P, \quad \text{Var}(\chi_{\nu_P}^2) = 2\nu_P. \quad (4.3)$$

The scaling in Equation (4.2) therefore matches the mean automatically:

$$\mathbb{E} \left[\mathbb{E}\{\widehat{V}(P)\} \frac{\chi_{\nu_P}^2}{\nu_P} \right] = \mathbb{E}\{\widehat{V}(P)\}. \quad (4.4)$$

Its variance is

$$\text{Var} \left[\mathbb{E}\{\widehat{V}(P)\} \frac{\chi_{\nu_P}^2}{\nu_P} \right] = \frac{[\mathbb{E}\{\widehat{V}(P)\}]^2}{\nu_P^2} \text{Var}(\chi_{\nu_P}^2) \quad (4.5)$$

$$= \frac{2[\mathbb{E}\{\widehat{V}(P)\}]^2}{\nu_P}. \quad (4.6)$$

Matching this quantity to $\text{Var}\{\widehat{V}(P)\}$ gives

$$\boxed{\nu_P = \frac{2[\mathbb{E}\{\widehat{V}(P)\}]^2}{\text{Var}\{\widehat{V}(P)\}}}. \quad (4.7)$$

To first order, $\mathbb{E}\{\widehat{V}(P)\} \approx V(P)$. The required calculation is therefore the sampling variance of $\widehat{V}(P)$.

4.2 How the MI Variance Changes with the Table

To find how much the estimated MI variance changes between samples, we first ask how the population variance responds when a small amount of probability is moved into one cell (Hampel, 1974; van der Vaart, 1998). Fix a cell (x, y) and define the perturbed distribution

$$P_\varepsilon = (1 - \varepsilon)P + \varepsilon\delta_{(x,y)}, \quad (4.8)$$

where $\delta_{(x,y)}$ places probability one on cell (x, y) . At $\varepsilon = 0$, the distribution is P . Increasing ε moves a small amount of probability toward the selected cell while preserving total probability one.

For an arbitrary cell (i, j) ,

$$p_{ij}(\varepsilon) = (1 - \varepsilon)p_{ij} + \varepsilon \mathbf{1}\{i = x, j = y\}. \quad (4.9)$$

The joint and marginal probability derivatives at the original distribution are

$$\left. \frac{dp_{ij}(\varepsilon)}{d\varepsilon} \right|_0 = \mathbf{1}\{i = x, j = y\} - p_{ij}, \quad (4.10)$$

$$\left. \frac{dp_{i+}(\varepsilon)}{d\varepsilon} \right|_0 = \mathbf{1}\{i = x\} - p_{i+}, \quad (4.11)$$

$$\left. \frac{dp_{+j}(\varepsilon)}{d\varepsilon} \right|_0 = \mathbf{1}\{j = y\} - p_{+j}. \quad (4.12)$$

The derivatives of the joint probabilities sum to zero because

$$\sum_{i,j} \{\mathbf{1}(i = x, j = y) - p_{ij}\} = 1 - 1 = 0. \quad (4.13)$$

Effect on pointwise MI and MI.

Along the perturbation path, pointwise MI is

$$\ell_{P_\varepsilon}(i, j) = \log p_{ij}(\varepsilon) - \log p_{i+}(\varepsilon) - \log p_{+j}(\varepsilon). \quad (4.14)$$

Differentiating the three logarithms and substituting the probability derivatives gives

$$\left. \frac{d}{d\varepsilon} \ell_{P_\varepsilon}(i, j) \right|_0 = \frac{\mathbf{1}\{i = x, j = y\}}{p_{ij}} - \frac{\mathbf{1}\{i = x\}}{p_{i+}} - \frac{\mathbf{1}\{j = y\}}{p_{+j}} + 1. \quad (4.15)$$

MI along the path is

$$I(P_\varepsilon) = \sum_{i,j} p_{ij}(\varepsilon) \ell_{P_\varepsilon}(i, j). \quad (4.16)$$

The product rule separates its derivative into the change in probability weights and the change in pointwise-MI values:

$$\begin{aligned} \left. \frac{dI(P_\varepsilon)}{d\varepsilon} \right|_0 &= \sum_{i,j} \{ \mathbf{1}(i=x, j=y) - p_{ij} \} \ell_P(i, j) \\ &\quad + \sum_{i,j} p_{ij} \left. \frac{d\ell_{P_\varepsilon}(i, j)}{d\varepsilon} \right|_0. \end{aligned} \quad (4.17)$$

The first sum is $\ell_P(x, y) - I(P)$. In the second sum, the selected-cell term contributes $p_{xy}/p_{xy} = 1$. The row term contributes $-\sum_j p_{xj}/p_{x+} = -1$; the column term likewise contributes -1 ; and the constant contributes $\sum_{i,j} p_{ij} = 1$. These four contributions cancel. Hence

$$\boxed{\left. \frac{dI(P_\varepsilon)}{d\varepsilon} \right|_0 = \ell_P(x, y) - I(P)}. \quad (4.18)$$

This result recovers the centred pointwise-MI influence function used in Chapter 3.

Effect on the MI variance.

The target is the first-order change in

$$V(P) = \text{Var}_P\{\ell_P(X, Y)\}. \quad (4.19)$$

Write this variance in second-moment form:

$$V(P) = \text{E}_P\{\ell_P(X, Y)^2\} - I(P)^2. \quad (4.20)$$

Define the observation-level variance sensitivity by

$$g_P(x, y) = \left. \frac{dV(P_\varepsilon)}{d\varepsilon} \right|_0. \quad (4.21)$$

Differentiating Equation (4.20) gives

$$g_P(x, y) = \left. \frac{d}{d\varepsilon} \text{E}_{P_\varepsilon}\{\ell_{P_\varepsilon}(X, Y)^2\} \right|_0 - 2I(P) \left. \frac{dI(P_\varepsilon)}{d\varepsilon} \right|_0. \quad (4.22)$$

It remains to differentiate

$$\text{E}_{P_\varepsilon}\{\ell_{P_\varepsilon}(X, Y)^2\} = \sum_{i,j} p_{ij}(\varepsilon) \ell_{P_\varepsilon}(i, j)^2. \quad (4.23)$$

The product rule gives

$$\begin{aligned} \left. \frac{d}{d\varepsilon} E_{P_\varepsilon} \{ \ell_{P_\varepsilon}(X, Y)^2 \} \right|_0 &= \sum_{i,j} \left. \frac{dp_{ij}(\varepsilon)}{d\varepsilon} \right|_0 \ell_P(i, j)^2 \\ &\quad + 2 \sum_{i,j} p_{ij} \ell_P(i, j) \left. \frac{d\ell_{P_\varepsilon}(i, j)}{d\varepsilon} \right|_0. \end{aligned} \quad (4.24)$$

The first sum is

$$\ell_P(x, y)^2 - E_P \{ \ell_P(X, Y)^2 \}. \quad (4.25)$$

For the second sum, define the probability-weighted row and column means

$$E_P \{ \ell_P(X, Y) \mid X = x \} = \frac{\sum_j p_{xj} \ell_P(x, j)}{p_{x+}}, \quad (4.26)$$

$$E_P \{ \ell_P(X, Y) \mid Y = y \} = \frac{\sum_i p_{iy} \ell_P(i, y)}{p_{+y}}. \quad (4.27)$$

Substitution of Equation (4.15) gives

$$\begin{aligned} &\sum_{i,j} p_{ij} \ell_P(i, j) \left. \frac{d\ell_{P_\varepsilon}(i, j)}{d\varepsilon} \right|_0 \\ &= \ell_P(x, y) - \frac{\sum_j p_{xj} \ell_P(x, j)}{p_{x+}} - \frac{\sum_i p_{iy} \ell_P(i, y)}{p_{+y}} + I(P) \\ &= \ell_P(x, y) - E_P \{ \ell_P(X, Y) \mid X = x \} - E_P \{ \ell_P(X, Y) \mid Y = y \} + I(P). \end{aligned} \quad (4.28)$$

Combining Equations (4.18), (4.22), (4.25), and (4.28) yields

$$\boxed{\begin{aligned} g_P(x, y) &= \{ \ell_P(x, y) - I(P) \}^2 - V(P) \\ &\quad + 2 \left[\ell_P(x, y) - E_P \{ \ell_P(X, Y) \mid X = x \} \right. \\ &\quad \left. - E_P \{ \ell_P(X, Y) \mid Y = y \} + I(P) \right]. \end{aligned}} \quad (4.29)$$

The first line is the change if the pointwise-MI values were held fixed. The bracket accounts for changes in those values, including the effects of the selected joint probability and both margins.

4.3 Component degrees of freedom and final test

Population component degrees of freedom.

Define the population variability of the observation-level sensitivities by

$$\tau^2(P) = \text{Var}_P\{g_P(X, Y)\}. \quad (4.30)$$

The sensitivities average to zero under P , as shown in Appendix A. Their variation, not their average, determines how much the estimated MI variance changes between samples.

Assumption 6 (Regular variance sensitivity). *For the first-order Expanded Welch construction, $\tau^2(P) > 0$ and $\tau^2(Q) > 0$.*

An observation in cell $Z_a^{(P)}$ can be represented by a table that places probability one in that cell, denoted $\delta_{Z_a^{(P)}}$. Relative to P , that observation contributes the table change $\delta_{Z_a^{(P)}} - P$. The empirical table averages these one-observation tables, so its difference from P is

$$\hat{P} - P = \frac{1}{n_P} \sum_{a=1}^{n_P} \{\delta_{Z_a^{(P)}} - P\}. \quad (4.31)$$

The plug-in $\hat{V}(P)$ applies the variance calculation to this empirical table. Equation (4.21) gives its first-order change for one observation's direction. Because a first-order change is linear in the table error (van der Vaart, 1998, Chapter 20), the changes from the observations can be averaged:

$$\hat{V}(P) - V(P) \approx \frac{1}{n_P} \sum_{a=1}^{n_P} g_P(Z_a^{(P)}). \quad (4.32)$$

The observations are independent, so

$$\text{Var}\{\hat{V}(P)\} \approx \frac{1}{n_P^2} \sum_{a=1}^{n_P} \text{Var}_P\{g_P(Z_a^{(P)})\} \quad (4.33)$$

$$= \frac{\tau^2(P)}{n_P}. \quad (4.34)$$

Substitution into Equation (4.7) gives the population component degree of freedom

$$\boxed{\nu_P \approx \frac{2n_P V(P)^2}{\tau^2(P)}}. \quad (4.35)$$

Equation (4.35) depends on relative, not absolute, uncertainty. Since

$$\frac{\text{Var}\{\hat{V}(P)\}}{V(P)^2} \approx \frac{\tau^2(P)}{n_P V(P)^2}, \quad (4.36)$$

the component degrees of freedom are approximately twice the reciprocal of this relative variance. A stable estimate relative to the size of $V(P)$ has large degrees of freedom; a relatively unstable estimate has smaller degrees of freedom and produces a heavier-tailed final reference. This first-order component formula assumes $\tau^2(P) > 0$. Positive $V(P)$ alone does not guarantee that condition. A small change in the table can have no first-order effect on $V(P)$, even when MI and $V(P)$ are positive. In that case, a higher-order analysis of $\hat{V}(P)$ would be needed to describe its finite-sample distribution.

Table-based component degrees of freedom.

The observed table replaces every population quantity in Equation (4.29) by its empirical counterpart. The sums below are over occupied cells, because pointwise MI is undefined at an empirical zero. For each occupied cell,

$$\begin{aligned} \hat{g}_P(i, j) = & \{\hat{\ell}_P(i, j) - \hat{I}(P)\}^2 - \hat{V}(P) \\ & + 2 \left[\hat{\ell}_P(i, j) - \frac{\sum_{j'} \hat{p}_{ij'} \hat{\ell}_P(i, j')}{\hat{p}_{i+}} - \frac{\sum_{i'} \hat{p}_{i'j} \hat{\ell}_P(i', j)}{\hat{p}_{+j}} + \hat{I}(P) \right]. \end{aligned} \quad (4.37)$$

In all weighted sums, an empty cell contributes zero; no logarithm or sensitivity needs to be evaluated there. The empirical variance is centred explicitly:

$$\hat{\tau}^2(P) = \sum_{i,j} \hat{p}_{ij} \left\{ \hat{g}_P(i, j) - \sum_{k,l} \hat{p}_{kl} \hat{g}_P(k, l) \right\}^2. \quad (4.38)$$

The estimated component degree of freedom is

$$\boxed{\hat{\nu}_P = \frac{2n_P \hat{V}(P)^2}{\hat{\tau}^2(P)}}. \quad (4.39)$$

Repeating the same calculation for the second table gives $\hat{\nu}_Q$.

Combining the two components.

Substituting the two variance contributions and component degrees of freedom

into Equation (4.1) gives

$$\hat{\nu}_{\text{expanded}} = \frac{\left\{ \hat{V}(P)/n_P + \hat{V}(Q)/n_Q \right\}^2}{\left\{ \hat{V}(P)/n_P \right\}^2 / \hat{\nu}_P + \left\{ \hat{V}(Q)/n_Q \right\}^2 / \hat{\nu}_Q}. \quad (4.40)$$

The expanded two-sided p-value is

$$p_{\text{expanded}} = 2 \left\{ 1 - F_{t_{\hat{\nu}_{\text{expanded}}}}(|T|) \right\}, \quad (4.41)$$

where F_{t_ν} denotes the cumulative distribution function of a Student t variable with ν degrees of freedom.

The expanded calculation does not change $\hat{\Delta}_{\text{BC}}$, $\widehat{\text{SE}}$, or T . It changes only the reference distribution used to evaluate that statistic.

For a fixed finite T , the Student t distribution gives a two-sided p-value at least as large as the standard normal distribution used by Wald. Under the implemented validity rules, every valid Expanded Welch result is also Wald-valid. Thus every Expanded Welch rejection is also a Wald rejection, including when invalid outputs count as non-rejections in the reported rates. The correction cannot raise rejection where Wald is already too conservative. Appendix A proves this ordering and examines a zero-component case in the combined degrees-of-freedom formula.

4.4 Worked example and implementation

Consider the two observed 3×3 tables

$$N^P = \begin{pmatrix} 35 & 10 & 5 \\ 10 & 20 & 5 \\ 5 & 5 & 10 \end{pmatrix}, \quad N^Q = \begin{pmatrix} 30 & 5 & 5 \\ 5 & 30 & 5 \\ 5 & 5 & 10 \end{pmatrix}. \quad (4.42)$$

Their totals are $n_P = 105$ and $n_Q = 100$. The plug-in MI estimates are 0.13693 and 0.25848 nats. Since $(3 - 1)(3 - 1) = 4$, the corrected difference is

$$\hat{\Delta}_{\text{BC}} = \left(0.13693 - \frac{4}{2(105)} \right) - \left(0.25848 - \frac{4}{2(100)} \right) = -0.12059. \quad (4.43)$$

The empirical influence variances are $\hat{V}(P) = 0.27189$ and $\hat{V}(Q) = 0.43371$.

Hence

$$\widehat{\text{SE}} = \left(\frac{0.27189}{105} + \frac{0.43371}{100} \right)^{1/2} = 0.08323, \quad T = -1.44898. \quad (4.44)$$

Normal Wald gives $p = 0.14734$.

For Expanded Welch, the table-based sensitivity calculation first gives $\hat{\tau}^2(P) = 0.93212$ and $\hat{\tau}^2(Q) = 0.82035$. Substitution into Equation (4.39) gives component degrees of freedom

$$\hat{\nu}_P = 16.65, \quad \hat{\nu}_Q = 45.86. \quad (4.45)$$

Equation (4.40) combines them into $\hat{\nu}_{\text{expanded}} = 59.03$, giving $p = 0.15264$. Both methods fail to reject at the 5% level for these tables. The example shows that Expanded Welch changes the reference, not the observed statistic.

Operational algorithm and complexity.

Table 4.1 summarises one comparison.

Table 4.1: Operational Expanded Welch–Satterthwaite algorithm.

Step	Calculation
1	Validate the two aligned count tables and their declared alphabet dimensions; independence of samples is a design assumption.
2	Compute empirical probabilities, pointwise MI, plug-in MI, and the leading bias-corrected difference.
3	Compute $\hat{V}(P)$, $\hat{V}(Q)$, the shared standard error, and T .
4	Compute the row and column pointwise-MI means needed in $\hat{g}_P$ and $\hat{g}_Q$.
5	Compute $\hat{\tau}^2(P)$, $\hat{\tau}^2(Q)$, and the two component degrees of freedom.
6	Combine the components using Equation (4.40).
7	Return the p-value from the Student t distribution, effective degrees of freedom, and validity diagnostics.

Each empirical probability, pointwise-MI value, row mean, column mean, and variance contribution is calculated by a fixed number of table scans. The time complexity is therefore $O(rc)$, and the working memory is also $O(rc)$. The method is deterministic once the two count tables are given.

4.5 Behaviour and limitations

Near independence.

The adjustment has a local explanation. Consider a population table that approaches independence while its row and column margins stay fixed. At indepen-

dence, MI is zero and its first-order change is also zero. The leading changes in MI, the pointwise-MI variance $V(P)$, and the variance sensitivity $\tau^2(P)$ can therefore be compared using a second-order expansion. Appendix A.5 gives the calculation and its fixed-margin assumptions. Close to independence on that path, it gives

$$V(P) \approx 2I(P), \quad \tau^2(P) \approx 4V(P). \quad (4.46)$$

Both relationships come from the same changes to the table; they are local approximations, not identities for all populations. Substituting them into the component degrees-of-freedom formula gives

$$\nu_P = \frac{2n_P V(P)^2}{\tau^2(P)} \approx \frac{n_P V(P)}{2} \approx n_P I(P). \quad (4.47)$$

The cancellation explains why a low starting MI can produce few effective degrees of freedom and a higher rejection cutoff from the Student t distribution. This result assumes fixed table size, positive cell probabilities and a population close to independence along the stated path; applying it to samples also requires the first-order approximation to be useful. It does not describe exact independence, where quadratic sampling variation matters and the first-order test may be poorly calibrated. Appendix E treats that boundary separately.

Theoretical status and validity conditions.

Equations (4.18) and (4.29) describe how MI and its estimated variance change when the underlying table changes slightly. For tables with positive cell probabilities, these changes give their first-order large-sample approximations.

Approximating $\hat{V}(P)$ by the scaled chi-squared model in Equation (4.2) and evaluating T against a Student t distribution are working approximations, not exact finite-sample laws. In particular, $\hat{I}(P)$ and $\hat{V}(P)$ come from the same sample and can vary together. Adjusting the reference for uncertainty in $\hat{V}(P)$ does not account fully for this dependence or make the statistic follow a Student t distribution exactly. Appendix A gives the first-order covariance. As sample sizes grow under regular conditions, the estimated degrees of freedom increase and the Student t distribution approaches the standard normal distribution used by Wald.

Both tests need a finite corrected MI difference and a positive, finite estimated standard error. Expanded Welch also needs positive, finite degrees of freedom for each population and for their combination. If a requirement fails, the calculation is recorded as invalid, not as evidence for the null. Appendix C.2 gives the numerical validity criteria used in the experiment. Chapters 5 and 6 report how often valid results are obtained and whether the correction improves calibration.

Chapter 5

Experimental Design

The purpose of the experiment is to compare the calibration, detection and validity of two tests across controlled populations. Each design point fixes two population distributions in advance. Repeated independent samples from those populations show how often each procedure returns a valid result and how often it rejects. Because the population MI values are known, a rejection at equal MI is a false positive; a rejection at unequal MI is detection.

5.1 A Concrete Pair of Population Tables

Consider two 2×2 populations whose row and column margins are $(0.7, 0.3)$ for P and $(0.8, 0.2)$ for Q . The construction gives

$$P \simeq \begin{pmatrix} 0.53273 & 0.16727 \\ 0.16727 & 0.13273 \end{pmatrix}, \quad Q \simeq \begin{pmatrix} 0.68175 & 0.11825 \\ 0.11825 & 0.08175 \end{pmatrix}. \quad (5.1)$$

In both populations, X and Y are dependent, with MI values of 0.02 and 0.03 nats, respectively. Their absolute difference, 0.01 nats, is the horizontal-axis value used in the figures.

5.2 Constructing fixed joint distributions

For one population, let $a = (a_1, \dots, a_r)$ and $b = (b_1, \dots, b_c)$ be prescribed row and column probabilities. The independence table has cell probabilities $a_i b_j$. Association is introduced by

$$R(t) = ab^\top + tH, \quad H\mathbf{1}_c = 0, \quad \mathbf{1}_r^\top H = 0. \quad (5.2)$$

The zero row and column sums mean that changing t leaves the margins fixed. For the main construction, H is zero except in the first two rows and columns, where it is

$$H_{1:2,1:2} = \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix}. \quad (5.3)$$

Thus t is added to the two diagonal cells and removed from the other two cells. The changes cancel within each affected row and column. The same matrix specifies a 2×2 association and a localised change in a larger table; only its embedding differs. This follows the minimal four-violation sparse alternative used by [Berrett and Samworth \(2021, Section 3\(b\)\)](#). Here it is adapted to compare two fixed populations: their margins are prescribed separately and the perturbation strength is chosen to attain a specified MI and is not used directly as the horizontal-axis effect.

For the pair in Equation (5.1), the starting independence tables are

$$P_0 = \begin{pmatrix} 0.49 & 0.21 \\ 0.21 & 0.09 \end{pmatrix}, \quad Q_0 = \begin{pmatrix} 0.64 & 0.16 \\ 0.16 & 0.04 \end{pmatrix}. \quad (5.4)$$

Adding the block in Equation (5.3) with strengths $t_P \simeq 0.04273$ and $t_Q \simeq 0.04175$, respectively, gives the displayed P and Q . The strengths are chosen to reach their separate MI targets; the 0.01-nat difference is not a cell-probability increment.

The largest nonnegative strength before a cell with $H_{ij} < 0$ reaches zero is

$$t_{\max} = \min_{H_{ij} < 0} \frac{a_i b_j}{-H_{ij}}. \quad (5.5)$$

The construction solves for $0 \leq t < t_{\max}$ such that $I\{R(t)\}$ equals the requested target. Targets at or beyond the endpoint are not admitted. This endpoint is specific to the chosen path, not the largest possible MI over every table with those margins. Both constructed tables are checked for strictly positive cells, correct margins and target MI before any sampling. Appendix B gives the numerical details.

For a dimension of size k , the skewed margin used in the design is

$$m_k(\rho) = \left(\rho, \frac{1-\rho}{k-1}, \dots, \frac{1-\rho}{k-1} \right). \quad (5.6)$$

The three principal profiles are specified separately for the two populations in Table 5.1. For a rectangular $r \times c$ table, the formula is applied with $k = r$ to row margins and $k = c$ to column margins. At a zero MI difference, the same-margin profiles produce $P = Q$ under this construction; the different-skew profile has $I(P) = I(Q)$ but $P \neq Q$.

Table 5.1: Row and column margins used for each population in the principal design. “Uniform” means the vector with every entry equal to $1/k$.

Profile	P : rows and columns	Q : rows and columns
Uniform	uniform	uniform
Same skew	$m_k(0.8)$	$m_k(0.8)$
Different skew	$m_k(0.7)$	$m_k(0.8)$

In the usual effect direction, $I(P) = I_0$ and $I(Q) = I_0 + |\Delta|$. The paired solve chooses separate strengths for P and Q ; it does not add random noise to their true MI values. The reverse-direction imbalance experiment assigns the higher target to P . The statistical convention remains $\Delta = I(P) - I(Q)$, which is negative in the usual direction. Figures use the nonnegative magnitude $|\Delta|$.

Two focused checks ask whether results depend on where the association sits in the table. The “rare block” moves the four changed cells from the upper-left corner to rows and columns with smaller probabilities in skewed populations. The “spread” pattern changes cells throughout the table while keeping the row and column totals fixed. The log-linear comparison also spreads association through the table, using a different rule. In each case, the population is built to have the specified margins and target MI. Tables can therefore agree on margins and MI while assigning different probabilities to individual cells. Appendix B gives the three construction formulas.

5.3 Sampling, tests and measured outcomes

For each unique configuration, 20,000 independent pairs of multinomial count tables are drawn from its fixed P and Q . Normal Wald and Expanded Welch are applied to the *same* count-table pair. This pairing makes method differences more precisely measurable than unrelated simulations, though it does not remove Monte Carlo uncertainty. The two-sided nominal threshold is $\alpha = 0.05$ throughout the statistical study.

For each method, the primary plotted rejection rate is

$$\frac{\#\{\text{sample pairs giving a valid p-value no larger than } 0.05\}}{20,000}. \quad (5.7)$$

At $|\Delta| = 0$ this estimates the false-positive rate. At $|\Delta| > 0$ it estimates the probability of detecting the specified difference at the fixed nominal threshold, often called power. A rate near 0.05 is desirable under the null; a smaller rate is not automatically better. Under an alternative, higher detection is useful only alongside

acceptable null calibration and a usable test result.

An invalid calculation contributes no rejection to Equation (5.7), but it does not count as evidence for the null. The valid-result rate shows how often each test returns a usable p-value. The reported 95% intervals show simulation uncertainty at each fixed population pair. Intervals for the difference between methods account for applying both tests to the same sampled tables. Appendix C gives the interval calculations and rates restricted to valid results.

The “minimum expected count” diagnostic is $\min_{ij} n_P p_{ij}$ for P , or $\min_{ij} n_Q q_{ij}$ for Q . It is not the expectation of the smallest observed count. No lower bound on that diagnostic was used to admit statistical configurations. Empty observed cells and rows remain possible. In the plots a hollow marker warns that fewer than 90% of outputs of that method were valid; 90% is a display convention, not a validated safety threshold.

5.4 Experiment families and computation

Table 5.2 records the question answered by each family. The common population construction above applies unless the named factor changes. Exact parameter sets are in Appendix B. The principal grid crosses square shapes $2 \times 2, 3 \times 3, 5 \times 5, 8 \times 8$ with the three margin profiles in Table 5.1. It fixes $I(P) = 0.02$ nats and uses $|\Delta| \in \{0, 0.001, 0.002, 0.005, 0.01, 0.02\}$ nats and equal sample sizes $n_P = n_Q \in \{2, 5, 10, 20, 50, 100, 250, 500, 1000\}$. For every feasible combination, the constructor fixes P and Q before the 20,000 table pairs are sampled and both tests are applied to each pair. The other families change one or more stated factors from this principal grid; Appendix B lists their exact values. There are 534 distinct population pairs (P, Q) , each specifying two fixed joint probability tables. Reusing a pair at different sample sizes creates distinct sampling configurations. There are 3,111 unique statistical configurations. Each unique configuration is sampled once, giving 62,220,000 table pairs in the final statistical run. The within-family counts in Table 5.2 do not sum to 3,111 because the same configuration can occur in several families. Results from the same configuration may therefore appear in more than one comparison; this is not an independent replication. Chapter 6 presents the main comparisons, with supporting figures in Appendix D.

Table 5.2: The fixed experiment families. All statistical points use 20,000 paired replicates and the 5% threshold.

Family	Configurations	Factor examined
Main	648	Square table size, three margin profiles and equal sample size.
Baseline	648	Starting positive MI of 0.0001, 0.001 or 0.02 nats.
Changed cells	144	First, rare or spread association patterns.
Imbalance	1,056	Sample allocation and which population has higher MI.
Larger effects	162	MI differences up to 0.2 nats where feasible.
Rectangular	108	2×3 and 3×5 tables.
Construction	432	Additive versus ordinal log-linear population tables.
Extreme skew	315	Dominant margins up to 0.9995, including $n = 1$ stress.
Rare-cell stress	140	First- versus rare-block changes, including $n = 1$ stress.
Independence	144	Exact $I(P) = 0$ boundary.
Convergence	108	Fixed equal-MI populations with sample sizes up to 50,000.
Runtime	60	Full-call timing regimes, separate from simulation.

The runtime experiment measures a full method call after count tables are already available. It varies table shape, profile, sample size and null or alternative status. Each regime uses 200 pre-generated table pairs, 20 warm-up calls and five timed passes. Thus it measures method computation, not the cost of constructing populations or generating random tables. Median call times and interquartile ranges are reported in Chapter 6.

Chapter 6

Experimental Results and Discussion

The figures plot how often each test rejects across 20,000 sampled table pairs. A point at a true MI difference of zero measures false positives; points to its right measure detection when the population MIs differ. The horizontal coordinate is fixed by the population construction, not estimated from the random samples. Blue solid circles show Normal Wald and magenta dashed squares show Expanded Welch. The dotted line marks the target 0.05 null rejection rate. Every plotted rate divides valid rejections by all 20,000 replicates; a hollow marker warns that fewer than 90% of calculations were valid. Where a zoom is shown, the upper panels use a 0–1 rejection axis and the lower panels enlarge the 0–0.15 range.

The comparisons proceed from null calibration to detection, sparse-sample validity, possible causes of error and larger-sample behaviour. Additional matched figures are in [Appendix D](#).

6.1 Calibration across table size and margins

The first comparison keeps both populations uniform in a 2×2 table. [Figure 6.1](#) shows three equal sample sizes and asks whether the lower Expanded Welch rejection rate is closer to the target 0.05 at each size.

For these 2×2 tables, both P and Q have uniform row and column margins $(1/2, 1/2)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in [Equation \(5.3\)](#), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure 6.1: Normal Wald and Expanded Welch rejection rates for 2×2 tables under the uniform margin profile. The lower panels enlarge the 0–0.15 range.

At zero difference, both methods reject below 0.05 at $n = 100$ and 1,000; Expanded Welch is lower still. At $n = 10$, Wald is slightly above 0.05 and Expanded Welch below it. Validity is high in all three panels. Rejection grows most clearly with MI difference at $n = 1,000$, where the two curves are close. At the smaller sizes, low rejection should not be mistaken for better detection.

Figure 6.2 asks whether the same comparison changes in a larger, still-uniform 8×8 table.

For these 8×8 tables, both P and Q have uniform row and column margins $(1/8, \dots, 1/8)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure 6.2: Normal Wald and Expanded Welch rejection rates for 8×8 tables under the uniform margin profile. The lower panels enlarge the 0–0.15 range.

At $n = 10$ and 100 , the zero-difference rates exceed 0.05 for both methods; Expanded Welch reduces the excess but does not always reach the target. These plotted calculations are valid, so the difference is not caused by missing outputs. At $n = 100$, rejection changes little over the displayed alternatives; at $n = 1,000$, it grows strongly. Even with uniform margins, table size and sample size alter the trade-off. The next comparison allows the two populations' margins to differ.

For these 8×8 tables, P has row and column margins $(0.7, 0.3/7, \dots, 0.3/7)$, while Q has $(0.8, 0.2/7, \dots, 0.2/7)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure 6.3: Normal Wald and Expanded Welch rejection rates for 8×8 tables under the different skew margin profile. The lower panels enlarge the 0–0.15 range.

At zero difference, both methods are too liberal at $n = 10$ and 100 , although Expanded Welch rejects less. Its $n = 10$ markers are hollow, so validity contributes to that reduction. At $n = 100$, rejection falls rather than grows over the displayed alternative range; at $n = 1,000$, detection rises but stays lower for Expanded Welch. This is the clearest reminder not to infer power from the null rate alone.

Across the 24 main null regimes with $n = 10$ or 20 , Expanded Welch is numerically closer to 0.05 in 15, whereas Wald is closer in 9. At $n = 500$ or $1,000$, Wald is closer in all 24 main null regimes. The $n = 5$ regimes need separate treatment below because Expanded Welch never rejects there. Appendix D shows the remaining shape–margin plots.

Table 6.1 makes the two directions of the calibration trade-off concrete in null settings where both methods return a p-value on every sample. Brackets beside each rate give pointwise Wilson 95% intervals; the final column gives an approximate paired 95% interval for Wald minus Expanded Welch on the same samples.

Table 6.1: Rejection rates and paired differences in two fully valid null regimes.

Regime	Wald	Expanded	Wald minus Expanded
	Rate [95% interval]	Rate [95% interval]	Rate [95% interval]
8×8 , uniform, $n = 20$	0.0687 [0.0653, 0.0723]	0.0570 [0.0538, 0.0602]	0.0118 [0.0103, 0.0132]
2×2 , uniform, $n = 1000$	0.0445 [0.0417, 0.0474]	0.0387 [0.0361, 0.0414]	0.0059 [0.0048, 0.0069]

In the 8×8 row, Expanded Welch moves Wald’s above-target rate closer to 0.05 . In the 2×2 row, it moves an already-low rate farther below the target. Invalid outputs cannot explain either difference. The reason for this one-sided trade-off is built into the methods: on the same valid table pair, Expanded Welch uses the same statistic as Wald but evaluates it against a Student t distribution, which gives a no-smaller two-sided p-value. It can lower rejection, not raise it. This makes detection under unequal MI the next question.

6.2 Detection as the true MI difference grows

The main grid shows only differences up to 0.02 nats. Figure 6.4 extends a uniform 2×2 comparison to 0.2 nats to check whether the low rejection over small effects persists as the true MI difference grows.

Both populations have row and column margins $(0.5, 0.5)$, use the first 2×2 change block, and have sample sizes $n_P = n_Q = 100$. The population MI values are $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The difference grid is $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02, 0.05, 0.1, 0.2$ nats.




Figure 6.4: Rejection rates over wider MI differences for uniform 2×2 tables.

At zero difference, both methods are below the 0.05 line; Expanded Welch is lower, and validity is high. Rejection then rises markedly as the true MI difference reaches the larger targets: by 0.2 nats both methods reject more than 90% of samples. Expanded Welch remains below Wald at intermediate differences, but the gap narrows when detection is already high. This wider axis shows that the low rejection over the main grid’s small differences is not a permanent ceiling on detection.

The next question is whether the starting MI changes this trade-off. Figure 6.5 compares three starting MI values in 3×3 different-skew tables while keeping the other displayed settings fixed.

From left to right, the columns use $I(P) = 0.0001, 0.001, 0.02$ nats, with $I(Q) = I(P) + |\Delta|$. The row and column margins are $(0.7, 0.15, 0.15)$ for P and $(0.8, 0.1, 0.1)$ for Q . Both populations use the first 2×2 change block and $n_P = n_Q = 100$, with $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure 6.5: Rejection rates at three starting MI values for 3×3 tables with different margins. The lower panels enlarge the 0–0.15 range.

For $I(P) = 0.0001, 0.001, 0.02$, the zero-difference Wald rates are 0.0193, 0.0146 and 0.0171; Expanded Welch gives 0.0017, 0.0016 and 0.0083. All these null calculations are valid. Rejection rises only modestly across the displayed alternatives, and both methods remain mostly below 0.05. Starting closer to independence does not make the higher cutoff from the Student t distribution helpful here: it lowers an already-conservative null rate and further reduces detection.

The local calculation in Chapter 4 explains why its degrees of freedom can become small near independence: the cutoff from the Student t distribution becomes more demanding where the first-order MI approximation is also less reliable. Alternative rejection is measured at the common nominal 5% threshold. Because the methods have different null rejection rates, the curves should not be read as a comparison at matched achieved false-positive rates. They show the practical trade-off under the procedures as implemented: in these examples, the extra caution reduces detection as well as false positives.

6.3 Validity under sparse sampling

An observed table can omit many cells even when every population cell has positive probability. Figure 6.6 pushes the margins toward one dominant category while holding $n_P = n_Q = 100$. It asks whether near-zero rejection in very sparse tables means good calibration or simply a failure to calculate a test.

From left to right, the dominant row and column probabilities for (P, Q) are $(0.9, 0.95)$, $(0.99, 0.995)$, and $(0.999, 0.9995)$. Each remaining margin entry divides the leftover probability equally between the two other categories. Both populations use the first 2×2 change block, with $I(P) = 0.00001$ and $I(Q) = I(P) + |\Delta|$. The difference grid is $|\Delta| = 0, 0.00001, 0.00005, 0.0001, 0.0002$ nats.




Figure 6.6: Rejection rates for 3×3 tables with increasingly dominant categories. The lower panels enlarge the 0–0.15 range.

For dominant margins 0.9 and 0.95, the null rates are 0.0701 for Wald and 0.0088 for Expanded Welch: the adjustment moves an excessive rate past 0.05. At the most extreme margins, 0.999 and 0.9995, both rates are zero, but only 0.0118 of Wald and 0.00005 of Expanded Welch null calculations are valid. Rejection does not meaningfully grow over the displayed alternatives there. The hollow markers show why near-zero unconditional rejection cannot be read as good calibration in sparse data.

Low validity explains much of the zero rejection in the most extreme panel. There is a second, distinct small-sample effect: Expanded Welch rejects in none of the 12 main-grid null regimes with $n = 5$, including uniform 3×3 , 5×5 and 8×8 regimes where most calculations are valid. Thus its zero rejection cannot always be explained by missing outputs. In one 8×8 different-skew alternative at $n = 5$, Wald rejects 0.30865 and Expanded Welch rejects 0.00000, while their valid-result rates are 0.83840 and 0.31415. Rejection and validity must be reported together.

Wald needs a positive usable total estimated variance. The expanded implementation additionally requires positive finite component sensitivities and degrees of

freedom. Some count tables therefore produce a Wald p-value but no Expanded Welch p-value. The primary rate counts such failures as non-rejections for the procedure; they are not evidence that the population MIs are equal. The exact numerical validity criteria are given in Appendix C.2.

6.4 Checks on population and sampling choices

The preceding figures isolate the central trade-off. Additional matched checks ask whether it depends on where the association sits, how the populations are constructed, the sample allocation or the table shape. Their plots and exact specifications are collected in Appendix D. Moving association between early, rare and spread cells changes the underlying table even when margins and MI are fixed (Figure D.10). An ordinal log-linear construction also changes rates under matched margins and MI (Figure D.14). Reversing which population has the larger sample or higher MI changes the alternative curves (Figures D.11 and D.12). Rectangular tables show that the conservative pattern is not exclusive to square shapes (Figure D.13). These checks limit how far any single plotted example can be generalised; none removes the need to judge null calibration and validity for the particular population pair.

The exact-independence check in Figure D.15 is kept with those supporting plots because it addresses a different mathematical boundary: the first-order MI variance becomes zero at independence. Appendix E explains why a one-table independence test has a different reference distribution.

6.5 Limits of changing the reference distribution

Both methods use the same estimated MI difference and the same plug-in standard error. Changing the reference distribution leaves errors in either quantity untouched. Even when the population MIs are equal, the two estimates can retain different small-sample biases if P and Q have different margins. Their leading equal-sample corrections cancel, but the remaining biases need not; the estimated difference can therefore be shifted away from zero. The plots do not show how much this contributes to excess rejection.

A supplementary check illustrates the shared-standard-error problem. For a 3×3 different-skew null with $I(P) = I(Q) = 0.02$ and $n = 100$, Wald rejects 0.0166 of independently simulated table pairs. Keeping the same estimated MI differences but dividing each by one fixed standard deviation estimated in a separate pilot

simulation raises rejection to 0.0551. The cutoff from the normal distribution has not changed. In this setting, the average squared plug-in standard error is about 17% larger than the measured variance of the MI difference. A larger standard error tends to make the absolute test statistic smaller and rejection less likely.

The fixed-denominator check cannot isolate that size effect: it also removes the way the estimated difference and standard error move together between samples. In the original Wald calculation, both come from the same observed tables. Appendix A.6 examines their first-order relationship. These calculations are diagnostics, not replacement tests. They show why changing only to a Student t distribution cannot correct every error in the statistic shared by Wald and Expanded Welch.

6.6 Larger samples and computational cost

Figure 6.7 keeps each population pair fixed under equal MI and increases the sample size. This tests whether the poor null rates are temporary small-sample effects for the chosen populations, and whether the answer changes with their common MI.

The rows compare 2×2 and 8×8 tables, and the columns use common population MI values of 0.0001 and 0.02 nats. For each shape, P 's row and column margins are $m_k(0.7)$ and Q 's are $m_k(0.8)$, with $k = 2$ or 8 . Both populations use the first 2×2 change block. The equal sample sizes are $n_P = n_Q = 1,000, 2,500, 10,000, 50,000$, plotted on a labelled logarithmic scale. Every point has equal population MI.




Figure 6.7: False-positive rates as sample size increases for 2×2 and 8×8 tables. The vertical axis ranges from 0 to 0.10.

All displayed calculations are valid. In the displayed 2×2 and 8×8 settings, rates move toward 0.05 more clearly when common MI is 0.02 than when it is 0.0001. Across all 12 first-block shape–margin combinations at $n = 50,000$, the 0.02-MI null rates range from 0.0470 to 0.0527 for Wald and 0.0466 to 0.0526 for Expanded Welch. At MI 0.0001, the ranges remain 0.0089–0.0280 and 0.0076–0.0138. Convergence can therefore be slow near independence.

These fixed populations do not establish a universal sample-size threshold for other tables. A separate consideration is computation. A full Expanded Welch method call has the same $O(rc)$ asymptotic complexity as Wald, but its extra sensitivity and degree-of-freedom calculations take longer in the measured implementation. Table 6.2 shows representative timings from the 60 timing regimes.

These examples use different-skew null populations with $n_P = n_Q = 100$ and $I(P) = I(Q) = 0.02$. Entries give the median and interquartile range, in microseconds, over five passes on 200 pre-generated table pairs per regime. The inputs are already available when timing begins.

Table 6.2: Median method-call times and interquartile ranges, in microseconds.

Shape	Normal Wald	Expanded Welch	Expanded/Wald
2×2	85.8 [85.5, 86.3]	150.0 [149.4, 151.3]	1.75
3×3	85.1 [84.6, 85.6]	150.0 [148.9, 150.9]	1.76
3×5	85.2 [84.9, 85.5]	150.1 [149.5, 150.9]	1.76
5×5	85.3 [85.1, 85.7]	149.2 [148.5, 150.1]	1.75
8×8	88.0 [87.1, 88.9]	153.4 [152.4, 154.8]	1.74

Across the 60 timing regimes, the median call time is about 85 microseconds for Wald and 150 for Expanded Welch, a ratio of about 1.75. These timings measure calls after tables are available; they exclude population construction and sampling. The extra cost is modest beside the statistical issue. Expanded Welch is useful where Wald is demonstrably too liberal and Expanded Welch remains valid, but the figures do not support replacing Wald across all regimes. Chapter 8 states the contribution and limits of that conclusion.

Chapter 7

Implications and Limitations

Chapter 6 interprets the figures as each question is answered. This chapter draws together what those results mean for using the test and how far the conclusions can be taken.

7.1 Implications for using Expanded Welch

Normal Wald and Expanded Welch use the same estimated MI difference and standard error. Expanded Welch can only make a valid two-sided p-value larger. This is helpful when Wald rejects an equal-MI null too often; it is harmful when Wald already rejects too rarely. Its extra component validity requirements can also reduce the number of samples for which a p-value is available. Lower rejection is therefore not by itself evidence of better calibration.

The tested populations include equal-MI pairs with different margins, so the comparison goes beyond the special case $P = Q$. They show that table size, margin shape, starting MI and sampling allocation can change the trade-off. No single numerical adjustment worked well across all those settings. For applications resembling a tested regime, Normal Wald remains the simpler baseline. Expanded Welch can be a conservative sensitivity check if a separate calibration study shows that Wald is too liberal and the expanded calculation is usually valid. Choosing the method after seeing which p-value is more favourable would not be a calibrated test.

The extra computation is not the principal obstacle. In the measured implementation, an Expanded Welch call takes about 1.75 times as long as a Wald call, but both have the same $O(rc)$ operation growth for an $r \times c$ table. The statistical limitations matter more than this timing difference.

7.2 Limits of the evidence

The study evaluates fixed, deliberately constructed population pairs, not a random sample from every possible joint distribution. Its 20,000 replicates per configuration make the rate for each fixed pair precise; they do not make the selected pairs representative of every application. The main construction changes four cells. Rare-cell, spread and log-linear constructions test other patterns, but the study does not pair equal-MI populations whose association is placed in different cells. That is a specific remaining design limitation.

The methods are studied with two independent samples from fixed categorical alphabets, a two-sided 5% threshold and the stated sample ladders. The analysis does not establish performance for paired observations, continuous MI estimators, growing alphabets or a real-data workflow. Exact independence is a separate boundary because the first-order MI variance vanishes there. The dedicated check and [Appendix E](#) describe that boundary but do not turn the positive-MI test into a general independence test.

The pointwise Monte Carlo intervals quantify simulation uncertainty for the chosen populations, not uncertainty about performance in untested populations.

These limits point to two distinct next steps: studying the joint finite-sample behaviour of the estimated MI difference and its standard error, and validating any alternative test under the equal-MI null when $P \neq Q$. A method for the stronger null $P = Q$ cannot be adopted for that question without checking its assumptions.

Chapter 8

Conclusion

This thesis studied how to compare the mutual information (MI) of two independently sampled categorical populations. Equal MI permits the population tables, their margins and their patterns of association to differ. The work makes two contributions to this problem. First, it derives how a change in one observation changes the plug-in estimate of MI variance, including the effects on cell probabilities and row and column totals. This gives an MI-specific Welch–Satterthwaite degrees-of-freedom calculation. Second, it evaluates that calculation against Normal Wald on fixed populations whose true MI differences are known, reporting false positives, detection, validity and runtime together.

The main finding is a qualified negative result. Expanded Welch reduces rejection relative to Wald and can move an excessive false-positive rate toward 5%. The same change makes an already-low rate lower and often reduces detection when MI differs. In sparse samples, it can also fail to return a p-value more often than Wald. The reference distribution alone cannot correct errors in the estimated MI difference or the standard error shared by both tests. This identifies both the value and the limit of a familiar variance-reference adjustment when applied to MI. Normal Wald remains the simpler general baseline for the tested task; Expanded Welch is best treated as a sensitivity check in a regime where its calibration and validity have been examined separately.

That recommendation is bounded by the experiment. Its population pairs were chosen to expose particular table sizes, margins, association patterns and sample sizes, not to represent all joint distributions. The derivation assumes independent observations, positive population cell probabilities and positive first-order MI variance. Exact independence requires a different approximation, and the study does not establish a rule for paired data, continuous MI estimators or growing alphabets.

The most useful next step is to study the estimated MI difference and its standard error together, especially near independence and in sparse tables. A likelihood-ratio or resampling method could provide a different comparison, but it would need validation for the actual null $I(P) = I(Q)$, including cases where $P \neq Q$. A real-data application and equal-MI pairs with association in different cells would then test how far the present simulation conclusions extend.

Appendix A

Additional Derivation Details

This appendix records algebra that supports Chapters 3 and 4 but is not required to follow the main argument.

A.1 Gradient of mutual information on the probability simplex

Mutual information can be written as

$$I(P) = \sum_{i,j} p_{ij} \log p_{ij} - \sum_i p_{i+} \log p_{i+} - \sum_j p_{+j} \log p_{+j}. \quad (\text{A.1})$$

Treating the cell probabilities as coordinates gives

$$\frac{\partial I(P)}{\partial p_{ij}} = (\log p_{ij} + 1) - (\log p_{i+} + 1) - (\log p_{+j} + 1) \quad (\text{A.2})$$

$$= \log \frac{p_{ij}}{p_{i+} p_{+j}} - 1 \quad (\text{A.3})$$

$$= \ell_P(i, j) - 1. \quad (\text{A.4})$$

Any valid probability perturbation $h = (h_{ij})$ satisfies $\sum_{i,j} h_{ij} = 0$. Its directional derivative is therefore

$$\left. \frac{d}{d\varepsilon} I(P + \varepsilon h) \right|_{\varepsilon=0} = \sum_{i,j} \{\ell_P(i, j) - 1\} h_{ij} \quad (\text{A.5})$$

$$= \sum_{i,j} \ell_P(i, j) h_{ij}. \quad (\text{A.6})$$

The constant in the unconstrained gradient disappears on the simplex. For the one-cell direction $h_{ij} = \mathbf{1}\{i = x, j = y\} - p_{ij}$, Equation (A.6) becomes

$$\ell_P(x, y) - \sum_{i,j} p_{ij} \ell_P(i, j) = \ell_P(x, y) - I(P), \quad (\text{A.7})$$

which is Equation (4.18).

A.2 Why the variance of pointwise MI appears

For empirical cell probabilities,

$$\hat{p}_{ij} - p_{ij} = \frac{1}{n_P} \sum_{a=1}^{n_P} \{\mathbf{1}(Z_a^{(P)} = (i, j)) - p_{ij}\}. \quad (\text{A.8})$$

Substitution into the first-order Taylor expansion of MI around P gives

$$\hat{I}(P) - I(P) \approx \sum_{i,j} \ell_P(i, j) (\hat{p}_{ij} - p_{ij}) \quad (\text{A.9})$$

$$= \frac{1}{n_P} \sum_{a=1}^{n_P} \left[\sum_{i,j} \ell_P(i, j) \{\mathbf{1}(Z_a^{(P)} = (i, j)) - p_{ij}\} \right] \quad (\text{A.10})$$

$$= \frac{1}{n_P} \sum_{a=1}^{n_P} \{\ell_P(Z_a^{(P)}) - I(P)\}. \quad (\text{A.11})$$

The observation-level summand has variance

$$\text{Var}_P\{\ell_P(X, Y) - I(P)\} = \mathbb{E}_P \left[\{\ell_P(X, Y) - I(P)\}^2 \right] \quad (\text{A.12})$$

$$= \sum_{i,j} p_{ij} \{\ell_P(i, j) - I(P)\}^2 \quad (\text{A.13})$$

$$= V(P). \quad (\text{A.14})$$

This is why the first-order variance of the plug-in MI estimator is $V(P)/n_P$, not a variance of the numerical category labels.

A.3 Mean of the variance influence function

The influence function $g_P(x, y)$ is centred at the population level. This can be shown without expanding its individual terms. The probability-weighted average of

the one-cell perturbation directions is

$$\sum_{x,y} p_{xy} \{\delta_{(x,y)} - P\} = P - P = 0. \quad (\text{A.15})$$

The first derivative is linear in the direction of change. Therefore, the weighted average of the one-cell derivatives is the derivative in the weighted-average direction:

$$E_P\{g_P(X, Y)\} = \sum_{x,y} p_{xy} \left. \frac{d}{d\varepsilon} V\{(1 - \varepsilon)P + \varepsilon\delta_{(x,y)}\} \right|_0 \quad (\text{A.16})$$

$$= \left. \frac{d}{d\varepsilon} V\left\{P + \varepsilon \sum_{x,y} p_{xy} \{\delta_{(x,y)} - P\}\right\} \right|_0 \quad (\text{A.17})$$

$$= \left. \frac{d}{d\varepsilon} V(P) \right|_0 = 0. \quad (\text{A.18})$$

The same weighted cancellation holds after exact plug-in substitution. The implementation nevertheless subtracts the probability-weighted empirical mean inside Equation (4.38) to protect the variance calculation from floating-point roundoff; this does not change the population definition.

A.4 Direct table substitutions

For observed counts $N^P = (N_{ij}^P)$, the quantities used by expanded Welch are calculated as

$$\hat{p}_{ij} = \frac{N_{ij}^P}{n_P}, \quad \hat{p}_{i+} = \sum_j \hat{p}_{ij}, \quad \hat{p}_{+j} = \sum_i \hat{p}_{ij}, \quad (\text{A.19})$$

$$\hat{\ell}_P(i, j) = \log \frac{\hat{p}_{ij}}{\hat{p}_{i+}\hat{p}_{+j}}, \quad \hat{I}(P) = \sum_{i,j} \hat{p}_{ij} \hat{\ell}_P(i, j), \quad (\text{A.20})$$

$$\hat{V}(P) = \sum_{i,j} \hat{p}_{ij} \{\hat{\ell}_P(i, j) - \hat{I}(P)\}^2. \quad (\text{A.21})$$

All sums involving pointwise MI are restricted to cells with $\hat{p}_{ij} > 0$. Direct substitution gives

$$\hat{g}_P(i, j) = \{\hat{\ell}_P(i, j) - \hat{I}(P)\}^2 - \hat{V}(P) \quad (\text{A.22})$$

$$+ 2 \left\{ \hat{\ell}_P(i, j) - \frac{\sum_{j'} \hat{p}_{ij'} \hat{\ell}_P(i, j')}{\hat{p}_{i+}} - \frac{\sum_{i'} \hat{p}_{i'j} \hat{\ell}_P(i', j)}{\hat{p}_{+j}} + \hat{I}(P) \right\}, \quad (\text{A.23})$$

$$\hat{\tau}^2(P) = \sum_{i,j} \hat{p}_{ij} \left\{ \hat{g}_P(i, j) - \sum_{k,l} \hat{p}_{kl} \hat{g}_P(k, l) \right\}^2, \quad (\text{A.24})$$

$$\hat{\nu}_P = \frac{2n_P \hat{V}(P)^2}{\hat{\tau}^2(P)}. \quad (\text{A.25})$$

These expressions are the direct $O(rc)$ implementation of the population derivation.

A.5 Fixed-margin behaviour near independence

Let $P_0 = (q_{ij})$ be independent, with positive cells $q_{ij} = a_i b_j$. Consider $P_t = P_0 + tH$, where every row and column sum of H is zero. For sufficiently small t , all cells remain positive and the margins remain a and b . We expand around $t = 0$, where $I(P_0) = 0$.

For $f(t) = I(P_t)$, the second-order Taylor formula is

$$f(t) = f(0) + f'(0)t + \frac{f''(0)}{2!}t^2 + O(t^3). \quad (\text{A.26})$$

The fixed margins make the pointwise MI in cell (i, j)

$$\ell_{P_t}(i, j) = \log \left(1 + \frac{tH_{ij}}{q_{ij}} \right) = \frac{tH_{ij}}{q_{ij}} - \frac{t^2 H_{ij}^2}{2q_{ij}^2} + O(t^3), \quad (\text{A.27})$$

using $\log(1 + z) = z - z^2/2 + O(z^3)$. Multiplying by the cell probability gives

$$(q_{ij} + tH_{ij})\ell_{P_t}(i, j) = tH_{ij} + \frac{t^2 H_{ij}^2}{2q_{ij}} + O(t^3). \quad (\text{A.28})$$

The first-order terms sum to zero because $\sum_{i,j} H_{ij} = 0$. Thus $f(0) = f'(0) = 0$, $f''(0) = \sum_{i,j} H_{ij}^2/q_{ij}$, and

$$I(P_t) = \frac{t^2}{2} \sum_{i,j} \frac{H_{ij}^2}{q_{ij}} + O(t^3). \quad (\text{A.29})$$

To calculate $V(P_t)$, square the leading pointwise-MI term. Its probability-weighted average is $t^2 \sum_{i,j} H_{ij}^2/q_{ij} + O(t^3)$, while $I(P_t)^2$ is of order t^4 . Hence

$$V(P_t) = t^2 \sum_{i,j} \frac{H_{ij}^2}{q_{ij}} + O(t^3). \quad (\text{A.30})$$

For the sensitivity in Equation (4.29), the row conditional mean of pointwise MI has first-order term $t \sum_j H_{ij}/a_i = 0$; the column mean likewise has first-order term $t \sum_i H_{ij}/b_j = 0$. The squared centred pointwise MI and $V(P_t)$ are both of order t^2 .

Therefore the bracket in Equation (4.29) supplies the first-order term:

$$g_{P_t}(i, j) = \frac{2tH_{ij}}{q_{ij}} + O(t^2). \quad (\text{A.31})$$

Since g_{P_t} has mean zero, its probability-weighted variance is

$$\tau^2(P_t) = 4t^2 \sum_{i,j} \frac{H_{ij}^2}{q_{ij}} + O(t^3). \quad (\text{A.32})$$

Comparing these three leading terms gives $V(P_t) \approx 2I(P_t)$ and $\tau^2(P_t) \approx 4V(P_t)$ for small positive MI along this path. At $t = 0$, the first-order sampling approximation itself degenerates; Appendix E treats that case separately.

A.6 Dependence between the MI and variance estimates

Both $\hat{I}(P)$ and $\hat{V}(P)$ are calculated from the same sample. Under fixed positive support, their first-order covariance is

$$\text{Cov}\{\hat{I}(P), \hat{V}(P)\} \approx \frac{1}{n_P} \mathbb{E}_P[\{\ell_P(X, Y) - I(P)\}g_P(X, Y)]. \quad (\text{A.33})$$

The corresponding expression for Q replaces P and n_P with Q and n_Q . To see its direction near independence, use the fixed-margin path $P_t = (q_{ij} + tH_{ij})$ from Section A.5. For small positive t , the expansions there give

$$\ell_{P_t}(i, j) - I(P_t) = \frac{tH_{ij}}{q_{ij}} + O(t^2), \quad g_{P_t}(i, j) = \frac{2tH_{ij}}{q_{ij}} + O(t^2). \quad (\text{A.34})$$

Substituting both expressions into the first-order covariance above gives

$$\text{Cov}\{\hat{I}(P_t), \hat{V}(P_t)\} \approx \frac{2t^2}{n_P} \sum_{i,j} \frac{H_{ij}^2}{q_{ij}} + O\left(\frac{t^3}{n_P}\right) \quad (\text{A.35})$$

$$= \frac{2V(P_t)}{n_P} + O\left(\frac{t^3}{n_P}\right). \quad (\text{A.36})$$

Thus the first-order large-sample approximation gives positive component covariance for small positive MI along this path. This is not a finite-sample guarantee or a result for exact independence: at $t = 0$, $V(P_0) = 0$ and the first-order MI approximation no longer describes its sampling variation. Nor does the component covariance alone determine the two-sample test's rejection rate.

A.7 Reference ordering and zero components

For positive component degrees of freedom, substituting Equation (4.39) and its Q analogue into Equation (4.40) gives the equivalent combined form

$$\hat{\nu}_{\text{expanded}} = \frac{\{\hat{V}(P)/n_P + \hat{V}(Q)/n_Q\}^2}{\hat{\tau}^2(P)/(2n_P^3) + \hat{\tau}^2(Q)/(2n_Q^3)}. \quad (\text{A.37})$$

This direct form remains defined when one table has $\hat{V} = \hat{\tau}^2 = 0$ and the other gives a positive denominator. The zero contribution drops out, leaving the other table's component degrees of freedom. The implementation nevertheless requires both components to have positive, finite degrees of freedom, so it records this case as invalid.

For a fixed statistic, the Student t distribution gives a p-value no smaller than Wald's. Write a variable with this distribution as $Z/\sqrt{S}$, where $Z \sim N(0, 1)$, $S \sim \chi_\nu^2/\nu$, and Z and S are independent. For $c > 0$, the conditional two-sided tail $2\{1 - \Phi(c\sqrt{s})\}$ is strictly convex in s , since its second derivative is $c\phi(c\sqrt{s})(s^{-3/2} + c^2s^{-1/2})/2 > 0$. As $\mathbb{E} S = 1$, Jensen's inequality gives $\Pr(|Z/\sqrt{S}| > c) > \Pr(|Z| > c)$; at $c = 0$ both tails equal one. Moreover, the implemented Expanded-valid set is a subset of the Wald-valid set. At a common fixed threshold, every Expanded rejection is therefore also a Wald rejection, both replicate by replicate and in unconditional rates that count invalid outputs as non-rejections.

A.8 Numerical derivative checks

The implementation tests the analytic derivatives against centred finite differences on positive probability tables. For a selected cell and a small positive step η , define

$$P_{\pm\eta} = (1 \mp \eta)P \pm \eta\delta_{(x,y)}. \quad (\text{A.38})$$

For sufficiently small η , the checks compare

$$\frac{I(P_\eta) - I(P_{-\eta})}{2\eta} \quad \text{with} \quad \ell_P(x, y) - I(P), \quad (\text{A.39})$$

$$\frac{V(P_\eta) - V(P_{-\eta})}{2\eta} \quad \text{with} \quad g_P(x, y). \quad (\text{A.40})$$

The perturbations preserve positive probabilities. These checks verify the algebra independently of the Monte Carlo experiment.

Appendix B

Population Construction and Exact Design Settings

The experiment uses fixed population tables. This appendix states the construction precisely enough to reproduce any plotted point and records the finite parameter grids.

B.1 Additive construction

For prescribed positive row vector a and column vector b , form $B = ab^\top$. The principal first-cell direction H contains $\begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix}$ in rows and columns 1 and 2 and zero elsewhere. The rare-cell direction places the same block in the last two rows and columns. The spread direction is

$$H_{ij} = u_i v_j, \quad u_i = -1 + \frac{2(i-1)}{r-1}, \quad v_j = -1 + \frac{2(j-1)}{c-1}. \quad (\text{B.1})$$

Both score vectors sum to zero, so every direction preserves margins.

For $t \geq 0$, define $R(t) = B + tH$. Calculate the first strength at which a decreasing cell reaches zero,

$$t_{\max} = \min_{H_{ij} < 0} \frac{B_{ij}}{-H_{ij}}. \quad (\text{B.2})$$

Evaluate the path endpoint MI. A target MI must lie strictly below this endpoint, within the stated numerical tolerance. For a positive target, Brent's one-dimensional root search ([Brent, 1973](#)) solves $I\{R(t)\} = I_{\text{target}}$ on $[0, t_{\max}]$; a zero target uses $t = 0$. The final table is accepted only if every cell is strictly positive, the sum is one, both margins agree with their prescribed vectors, and achieved MI agrees with

the target. The normalisation and margin tolerances are 10^{-12} ; the MI tolerance is $10^{-12} + 10^{-8}I_{\text{target}}$. These numerical checks ensure that a horizontal-axis difference denotes constructed population truth, not an empirical estimate. They do not impose an expected-count floor on the subsequent sample.

The row and column vectors are separately constructed for each population. For shape $r \times c$, the uniform profile uses uniform vectors of lengths r and c for both P and Q . The same-skew profile uses $m_r(0.8)$, $m_c(0.8)$ for both. The different-skew profile uses $m_r(0.7)$, $m_c(0.7)$ for P and $m_r(0.8)$, $m_c(0.8)$ for Q . The extreme regimes instead state the two dominant probabilities directly. The direction of the alternative is a separate parameter and should not be inferred from which population has the more skewed margin.

The log-linear comparison uses the scores u_i, v_j from Equation (B.1), but starts from positive cell weights

$$K_{ij}(\lambda) = \exp(\lambda u_i v_j) \quad (\text{B.3})$$

instead of adding tH to the independence table B . Iterative proportional fitting (Deming and Stephan, 1940) scales rows and columns of these weights until the prescribed margins are restored. At $\lambda = 0$, this gives the independence table; a numerical search then chooses λ to reach the target MI. The final table passes the same positivity, normalisation, margin and target-MI checks as the additive table. Unlike the first-block additive construction used in the comparison, it changes cells throughout the table. Matching margins and MI does not imply matching all cell probabilities or expected counts.

B.2 Statistical experiment grids

Unless a row of Table B.1 says otherwise, both populations use the first-cell additive direction, $I(P) = I_0$, $I(Q) = I_0 + |\Delta|$, and the three standard margin profiles. The symbol n means $n_P = n_Q = n$. Only targets producing strictly positive population tables are used. Curves should not be extrapolated beyond their last constructed target.

Table B.1: Exact statistical experiment grids; MI values are in nats.

Family	Settings
Main	Shapes $2 \times 2, 3 \times 3, 5 \times 5, 8 \times 8$; three profiles; $I_0 = 0.02$; $ \Delta \in \{0, 0.001, 0.002, 0.005, 0.01, 0.02\}$; $n \in \{2, 5, 10, 20, 50, 100, 250, 500, 1000\}$.

Family	Settings
Baseline	Main shapes and profiles; $I_0 \in \{0.0001, 0.001, 0.02\}$; main difference grid; $n \in \{20, 100, 1000\}$.
Changed cells	Shapes $3 \times 3, 5 \times 5, 8 \times 8$; different-skew margins; $I_0 = 0.0001$; $ \Delta \in \{0, 0.0001, 0.00025, 0.0005, 0.001, 0.002\}$; $n \in \{20, 100, 1000\}$. First and rare blocks for all shapes; spread pattern for 5×5 and 8×8 .
Imbalance	Main shapes, profiles and difference grid; $I_0 = 0.02$. Sample pairs are $(50, 50), (50, 100), (100, 50), (50, 250), (250, 50), (50, 500), (500, 50), (500, 500)$. Both Q -higher and P -higher MI directions.
Larger effects	Shapes $2 \times 2, 3 \times 3$; three profiles; $I_0 = 0.02$; main difference grid plus $0.05, 0.1, 0.2$; $n \in \{20, 100, 1000\}$.
Rectangular	Shapes $2 \times 3, 3 \times 5$; three profiles; $I_0 = 0.02$; main difference grid; $n \in \{20, 100, 1000\}$.
Construction	Shapes $3 \times 3, 5 \times 5, 8 \times 8$; uniform and different-skew profiles; $I_0 \in \{0.0001, 0.02\}$; main difference grid; $n \in \{20, 100, 1000\}$; additive first block and ordinal log-linear constructors.
Extreme skew	Shapes $2 \times 2, 3 \times 3, 8 \times 8$; dominant margin probability pairs $(\rho_P, \rho_Q) \in \{(0.9, 0.95), (0.99, 0.995), (0.999, 0.9995)\}$; $I_0 = 0.00001$; $ \Delta \in \{0, 0.00001, 0.00005, 0.0001, 0.0002\}$; $n \in \{1, 2, 5, 20, 100, 1000, 10000\}$.
Rare-cell stress	Shapes $3 \times 3, 8 \times 8$; dominant pair $(0.9, 0.95)$; first and rare blocks; $I_0 = 0.00001$; $ \Delta \in \{0, 0.00001, 0.000025, 0.00005, 0.0001\}$; same n ladder as extreme skew.
Independence	Shapes $2 \times 2, 8 \times 8$; three profiles; $I_0 = 0$; main difference grid; $n \in \{2, 20, 100, 1000\}$.
Convergence	Main square shapes and profiles; equal MI $I(P) = I(Q) \in \{0.0001, 0.02\}$; $n \in \{1000, 2500, 10000, 50000\}$. Additional rare-block controls use $3 \times 3, 5 \times 5, 8 \times 8$, different-skew margins and common MI 0.0001 .

Family	Settings
Runtime	Shapes $2 \times 2, 3 \times 3, 3 \times 5, 5 \times 5, 8 \times 8$; uniform and different-skew profiles; $I_0 = 0.02$; $ \Delta \in \{0, 0.02\}$; $n \in \{20, 100, 1000\}$. Each of the 60 regimes uses 200 pre-generated table pairs, 20 warm-up calls and five timed passes.

At $n = 1$, neither statistic is evaluable, so the result is recorded as invalid. Every unique statistical configuration has 20,000 sampled table pairs. When the same population pair and sample sizes appear in another section, the same simulation results are shown rather than generating a new set of samples.

Appendix C

Simulation Uncertainty and Validity Diagnostics

This appendix explains the rejection-rate denominators, uncertainty intervals and validity checks used in Chapters 5 and 6.

C.1 Rejection rates and simulation uncertainty

For each configuration, 20,000 table pairs are sampled. The plotted rate is the number of valid rejections divided by 20,000. The method-valid conditional rate divides valid rejections by that method’s valid outputs; it is undefined if there are none. The common-valid rate restricts to outputs where both methods are valid. A paired difference is based on counts of the four joint rejection outcomes for the same table pairs. These distinctions allow a reader to check whether a curve difference is primarily a changed reference or lost validity.

For the paired comparison, let D_s equal the Wald valid-rejection indicator minus the Expanded Welch valid-rejection indicator for sampled table pair s , matching the sign used in Table 6.1. Its average is $\bar{D} = (1/20,000) \sum_{s=1}^{20,000} D_s$, with approximate 95% interval

$$\bar{D} \pm 1.96 \left\{ \frac{1}{20,000(19,999)} \sum_{s=1}^{20,000} (D_s - \bar{D})^2 \right\}^{1/2}. \quad (\text{C.1})$$

This calculation preserves the within-replicate pairing. The individual method intervals use Wilson’s binomial score construction (Wilson, 1927).

C.2 Validity and sparse-sample diagnostics

Both tests require matching table shapes with at least two rows and columns, finite nonnegative integer counts, and more than one observation in each table. Normal Wald requires a finite corrected MI difference, a positive finite standard error, $\hat{V}(P) + \hat{V}(Q) > 10^{-14}$, and a finite p-value. The variance threshold treats values extremely close to zero as numerically unusable. Expanded Welch additionally requires positive finite degrees of freedom for both populations and their combination. If one component's degrees of freedom cannot be calculated, the expanded test is invalid even when the direct combined expression in Appendix A has a value. Failed calculations contribute no rejection to the unconditional rate.

For any P population with sample size n_P , its minimum expected cell count is $\min_{ij} n_P p_{ij}$; define Q 's analogously. No "expected minimum" of random observed counts is calculated. The sparsity diagnostics include empty observed rows and columns and the fraction of zero-count cells. Such values describe why a statistic might fail, but do not constitute a universal validity threshold.

The thesis does not treat a descriptive count of regimes closer to 0.05 as a formal paired test over regimes. Wilson intervals quantify binomial Monte Carlo variation at each fixed point, and paired normal intervals quantify a specified within-configuration method difference. Neither interval family is a simultaneous assertion across all configurations. Tables sharing a population but changing n or $|\Delta|$ answer different fixed-design questions and should not be collapsed into one mean rate.

C.3 Independent-pilot standard-error check

The diagnostic in Section 6.5 uses two independent sets of 20,000 sampled table pairs from the stated null populations. The pilot set gives the sample standard deviation of the corrected MI differences. The evaluation set is tested twice: once with each pair's plug-in standard error, and once with that fixed pilot standard deviation. Both calculations use the standard normal distribution and the two-sided 5% threshold. The comparison of variance estimates divides the evaluation set's average squared plug-in standard error by the squared pilot standard deviation. This check retains the population pair and estimated differences while changing how they are standardised.

Appendix D

Additional Experimental Figures

Chapter 6 uses seven figures to develop the main argument. The figures below retain the other matched comparisons from the same simulation run. Population and sample settings precede each figure, and its discussion follows in ordinary prose. Each plotted point uses 20,000 paired samples; hollow markers show validity below 90%. Except where stated, upper panels use a 0–1 rejection scale and lower panels enlarge 0–0.15.

D.1 Remaining square-table comparisons

The principal design crosses four square table sizes with three margin profiles. Chapter 6 discusses the uniform 2×2 and 8×8 comparisons and the different-skew 8×8 comparison. The nine other shape–margin combinations follow in the same order as the design grid.

For these 2×2 tables, both P and Q have row and column margins $m_2(0.8) = (0.8, 0.2)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.1: Normal Wald and Expanded Welch rejection rates for 2×2 tables under the same skew margin profile. The lower panels enlarge the 0–0.15 range.

The null points are below 0.05, especially for Expanded Welch. At $n = 10$, many Expanded Welch calculations are invalid (hollow markers), so its low unconditional rate has more than one cause. Detection rises with the constructed difference mainly at $n = 1,000$; there Wald remains slightly higher. This skewed pair gives no general calibration advantage to the more conservative reference.

For these 2×2 tables, P has row and column margins $m_2(0.7) = (0.7, 0.3)$, while Q has $m_2(0.8) = (0.8, 0.2)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.2: Normal Wald and Expanded Welch rejection rates for 2×2 tables under the different skew margin profile. The lower panels enlarge the 0–0.15 range.

At zero difference, the two tables have equal MI but different margins. Both methods are mostly below the 0.05 target, and Expanded Welch is lower. At $n = 10$, its hollow markers also indicate substantial invalidity. The alternative curves rise noticeably at $n = 1,000$, with Wald just above Expanded Welch. The same conservative pattern therefore persists when the equal-MI null does not require identical populations.

For these 3×3 tables, both P and Q have uniform row and column margins $(1/3, 1/3, 1/3)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.3: Normal Wald and Expanded Welch rejection rates for 3×3 tables under the uniform margin profile. The lower panels enlarge the 0–0.15 range.

At $n = 10$, Wald’s zero-difference rate is above 0.05 and Expanded Welch brings it below the line; both methods are almost always valid. At $n = 100$, however, both start well below 0.05 and the correction lowers rejection further. Detection grows strongly at $n = 1,000$, where Wald remains slightly higher. Whether the adjustment helps calibration changes with sample size, not just table shape.

For these 3×3 tables, both P and Q have row and column margins $m_3(0.8) = (0.8, 0.1, 0.1)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.4: Normal Wald and Expanded Welch rejection rates for 3×3 tables under the same skew margin profile. The lower panels enlarge the 0–0.15 range.

The $n = 10$ null rate falls from slightly above 0.05 for Wald to below it for Expanded Welch, but Expanded Welch is valid in only about two thirds of those samples. At $n = 100$ and 1,000, both null rates are below 0.05. Rejection rises with the MI difference most visibly at $n = 1,000$; Expanded Welch remains lower. The small-sample null reduction is not solely a reference-distribution effect.

For these 3×3 tables, P has row and column margins $m_3(0.7) = (0.7, 0.15, 0.15)$, while Q has $m_3(0.8) = (0.8, 0.1, 0.1)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.5: Normal Wald and Expanded Welch rejection rates for 3×3 tables under the different skew margin profile. The lower panels enlarge the 0–0.15 range.

With different margins, Wald rejects too often at the $n = 10$ null; Expanded Welch rejects less, although its hollow markers show reduced validity. At $n = 100$, both are conservative. At $n = 1,000$, detection increases with the true difference and the methods draw closer. This comparison shows a small-sample benefit against liberal rejection but not a consistent improvement across sample sizes.

For these 5×5 tables, both P and Q have uniform row and column margins $(1/5, \dots, 1/5)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.6: Normal Wald and Expanded Welch rejection rates for 5×5 tables under the uniform margin profile. The lower panels enlarge the 0–0.15 range.

Both methods are valid in the displayed panels. At zero difference, Wald is liberal at $n = 10$, while Expanded Welch reduces but does not eliminate the excess. At $n = 100$ and $1,000$, both null rates are below 0.05 . Little detection is gained over the displayed difference range at $n = 100$, whereas rejection grows clearly at $n = 1,000$. The correction lowers both false positives and detection.

For these 5×5 tables, both P and Q have row and column margins $m_5(0.8) = (0.8, 0.05, 0.05, 0.05, 0.05)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.7: Normal Wald and Expanded Welch rejection rates for 5×5 tables under the same skew margin profile. The lower panels enlarge the 0–0.15 range.

At $n = 10$, Wald starts above the 0.05 line and Expanded Welch below it, but the latter is invalid for roughly one third of samples. At $n = 100$ and 1,000, both begin below the line. Detection increases substantially only in the largest-sample panel, with Wald higher. The small-sample appearance of better calibration must be read together with validity.

For these 5×5 tables, the margins are $(0.7, 0.075, 0.075, 0.075, 0.075)$ for P and $(0.8, 0.05, 0.05, 0.05, 0.05)$ for Q , each used for both rows and columns. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.8: Normal Wald and Expanded Welch rejection rates for 5×5 tables under the different skew margin profile. The lower panels enlarge the 0–0.15 range.

The $n = 10$ null is liberal for both methods, though less so for Expanded Welch; its hollow markers also show reduced validity. At $n = 100$, rejection stays low and even falls across this short alternative range, so these curves should not be described as increasing power. At $n = 1,000$, detection rises and Wald remains slightly higher. The benefit of reducing excess null rejection is regime-specific.

For these 8×8 tables, both P and Q have row and column margins $m_8(0.8) = (0.8, 0.2/7, \dots, 0.2/7)$. From left to right, the columns use equal sample sizes $n_P = n_Q = 10, 100, 1000$. Both populations use the first 2×2 change block in Equation (5.3), with $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$. The MI difference varies over $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.9: Normal Wald and Expanded Welch rejection rates for 8×8 tables under the same skew margin profile. The lower panels enlarge the 0–0.15 range.

At $n = 10$, Wald is above 0.05 at the null and Expanded Welch below it, but many Expanded Welch outputs are invalid. Both null rates are below 0.05 at $n = 100$ and 1,000. Across the alternative points, appreciable growth appears chiefly at $n = 1,000$, with Wald higher. The adjustment does not turn the smaller-sample curves into strong detection.

D.2 Other controlled comparisons

The following figures hold the stated margins, MI targets and sample sizes fixed while changing association location, sample allocation, table shape or population construction. The final figure starts at exact independence, where the first-order MI variance is zero.

Figure D.10 compares the location of association in 5×5 tables. Columns place association in the first 2×2 block, last (rare) 2×2 block, or a cen-

tred score outer product spread across all cells. For every column, P 's row and column margins are $(0.7, 0.075, 0.075, 0.075, 0.075)$, and Q 's are $(0.8, 0.05, 0.05, 0.05, 0.05)$. Here $I(P) = 0.0001$, $I(Q) = I(P) + |\Delta|$, $n_P = n_Q = 100$, and $|\Delta| = 0, 0.0001, 0.00025, 0.0005, 0.001, 0.002$ nats.




Figure D.10: Rejection rates for 5×5 tables with association in different cells.

For first, rare and spread patterns, respectively, the zero-difference Wald rates are 0.0169, 0.0170 and 0.0160; Expanded Welch gives 0.0039, 0.0044 and 0.0041. All displayed calculations are valid. The alternative curves remain nearly flat over this short MI range, so neither method detects these differences well. Changing the location of association does not remove the conservative gap here, although this one plot does not make all patterns equivalent.

Figure D.11 compares unequal samples when Q has the higher MI. Columns have $(n_P, n_Q) = (50, 250)$ and $(250, 50)$. For both columns, P has 3×3 row and column margins $(0.7, 0.15, 0.15)$ and Q has $(0.8, 0.1, 0.1)$; the first-cell change block is used. $I(P) = 0.02$ and $I(Q) = 0.02 + |\Delta|$, with $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.11: Rejection rates under two sample allocations when Q has the higher MI.

At zero difference, Wald is near or above 0.05 while Expanded Welch is below it in both allocations. When Q has more observations $(n_P, n_Q) = (50, 250)$, rejection rises with $I(Q) - I(P)$; at 0.02 nats it reaches 0.0994 for Wald and 0.0247 for Expanded Welch. With P larger, the curves rise much less. These points are valid, but the alternative gap is not a size-adjusted power comparison because their matched null rates differ.

Figure D.12 reverses the MI direction. Columns have $(n_P, n_Q) = (50, 250)$ and $(250, 50)$. P 's 3×3 row and column margins are $(0.7, 0.15, 0.15)$; Q 's are $(0.8, 0.1, 0.1)$. Both use the first-cell change block. $I(Q) = 0.02$ and $I(P) = 0.02 + |\Delta|$, with $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.12: Rejection rates under two sample allocations when P has the higher MI.

The zero-difference rates are the same as in Figure D.11, because each allocation uses the same pair of populations in both figures at that point; P and Q still have different margins. Reversing the MI direction changes what happens next: when P has the larger sample, Wald rises well above its null rate, whereas it stays nearly flat when Q has the larger sample. At $(n_P, n_Q) = (50, 250)$ and 0.02 nats, Wald rejects 0.0495 and Expanded Welch 0.0232; both calculations are valid. Direction and allocation therefore matter even at the same absolute MI difference.

Figure D.13 compares rectangular tables. Columns are 2×3 and 3×5 . In each dimension k , P 's corresponding margin is $m_k(0.7)$ and Q 's is $m_k(0.8)$, for both rows and columns. Thus the first column uses P rows $(0.7, 0.3)$, columns $(0.7, 0.15, 0.15)$, and Q rows $(0.8, 0.2)$, columns $(0.8, 0.1, 0.1)$; the second applies the same formula to 3 and 5 categories. Both use first-cell changes, $n_P = n_Q = 100$, $I(P) = 0.02$, and $I(Q) = 0.02 + |\Delta|$ for $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.13: Rejection rates for 2×3 and 3×5 tables with different margins.

Both rectangular panels start below 0.05, with Expanded Welch lower. Rejection grows only modestly by 0.02 nats and remains below 0.05 for both methods throughout the displayed range; validity is high. Thus these shapes show weak detection over the chosen differences, not an improvement in false-positive calibration from the lower Expanded Welch curve. The conservative pattern is not confined to square tables.

Figure D.14 compares additive and ordinal log-linear constructions for 5×5 tables. The left column uses the first-cell additive block; the right fits an ordinal interaction while preserving the margins. In both, P 's row and column margins are $(0.7, 0.075, 0.075, 0.075, 0.075)$ and Q 's are $(0.8, 0.05, 0.05, 0.05, 0.05)$. Both have $I(P) = 0.02$, $I(Q) = 0.02 + |\Delta|$, $n_P = n_Q = 100$, and $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.14: Rejection rates under additive and ordinal log-linear constructions for 5×5 tables.

At zero difference, the additive and log-linear Wald rates are 0.0250 and 0.0181; Expanded Welch gives 0.0136 and 0.0119. All four are valid and below 0.05. Rejection stays low or falls slightly as the displayed difference grows. Matching margins and MI therefore does not make finite-sample behaviour identical across constructions, but neither construction removes the conservative gap between the methods.

Figure D.15 starts at exact independence in 2×2 tables. Columns are uniform margins $(0.5, 0.5)$ for both populations, same skew $(0.8, 0.2)$ for both, and different skew with $P = (0.7, 0.3)$ and $Q = (0.8, 0.2)$; each vector is used for both rows and columns. The first-cell block is used, $I(P) = 0$, $I(Q) = |\Delta|$, $n_P = n_Q = 100$, and $|\Delta| = 0, 0.001, 0.002, 0.005, 0.01, 0.02$ nats.




Figure D.15: Rejection rates for 2×2 tables starting at exact independence.

At the uniform zero-difference point, Wald rejects 0.00025 and Expanded Welch 0.00010 of samples; validity is high. Rejection grows little over most of the displayed alternative range, and Expanded Welch remains lower. This is not a routine positive-MI validation: at exact independence the population first-order MI variance is zero, so the approximation underlying both curves changes. Appendix E explains the different reference used by a one-table independence test.

D.3 Main-grid null-rate overview

Table D.1 describes where the 108 main-grid null rates fall relative to 0.025–0.075. This interval is only a descriptive aid, not a formal calibration test.

Each margin profile contains 36 configurations. The final column counts configurations with validity below 90% and overlaps the three rate-location columns. All four $n = 2$ configurations per profile have zero validity for both methods. Uniform and same-skew nulls have $P = Q$, while different-skew nulls have $P \neq Q$ and $I(P) = I(Q)$.

Table D.1: Main-grid null configurations below, inside or above the descriptive rejection-rate interval 0.025–0.075.

Margins	Method	Below	Inside	Above	Validity below 90%
Uniform	Normal Wald	10	18	8	4
Uniform	Expanded Welch	19	16	1	5
Same skew	Normal Wald	14	15	7	8
Same skew	Expanded Welch	25	11	0	12
Different skew	Normal Wald	9	16	11	8
Different skew	Expanded Welch	19	13	4	12

Appendix E

The Independence Boundary

The main thesis compares populations with $I(P) = I(Q) > 0$, allowing $P \neq Q$. This appendix explains why constructing an independent reference distribution does not turn the Expanded Welch method into a new first-order test of independence.

E.1 Constructing the independent reference

For a joint distribution $P = (p_{ij})$, define

$$q_{ij} = p_{i+}p_{+j}. \quad (\text{E.1})$$

Then $Q = (q_{ij})$ is a probability distribution because

$$\sum_{i,j} q_{ij} = \left(\sum_i p_{i+} \right) \left(\sum_j p_{+j} \right) = 1. \quad (\text{E.2})$$

It has the same margins as P , since

$$\sum_j q_{ij} = p_{i+}, \quad \sum_i q_{ij} = p_{+j}. \quad (\text{E.3})$$

The factorisation in Equation (E.1) makes X and Y independent under Q , so $I(Q) = 0$.

The divergence from P to this reference is

$$D_{\text{KL}}(P||Q) = \sum_{i,j} p_{ij} \log \frac{p_{ij}}{q_{ij}} \quad (\text{E.4})$$

$$= \sum_{i,j} p_{ij} \log \frac{p_{ij}}{p_{i+}p_{+j}} \quad (\text{E.5})$$

$$= I(P). \quad (\text{E.6})$$

Consequently, testing $I(P) = I(Q)$ for this constructed reference is exactly testing $I(P) = 0$, or independence.

E.2 Why the first-order term vanishes

At independence, $p_{ij} = p_{i+}p_{+j}$, and every population pointwise-MI value is zero:

$$\ell_P(i, j) = \log 1 = 0. \quad (\text{E.7})$$

Equation (4.18) therefore gives

$$\left. \frac{dI(P_\varepsilon)}{d\varepsilon} \right|_0 = \ell_P(x, y) - I(P) = 0. \quad (\text{E.8})$$

The first-order variance also vanishes:

$$V(P) = \text{Var}_P\{\ell_P(X, Y)\} = 0. \quad (\text{E.9})$$

The denominator used by Normal Wald and Expanded Welch is thus degenerate under the population independence null. This does not mean that a finite sample must have zero plug-in MI or zero plug-in variance. Sampling noise usually makes the empirical table depart from exact independence. The point is that these random departures do not enter MI through the usual first-order term at the null. They enter through the second-order term derived next, so changing only the degrees of freedom of a first-order statistic cannot supply the missing source of variation.

E.3 Second-order behaviour

Fix a strictly positive independent table P , and let $h = (h_{ij})$ describe a small redistribution of probability, with $\sum_{i,j} h_{ij} = 0$. Define the path

$$P_\varepsilon = P + \varepsilon h, \quad f(\varepsilon) = I(P_\varepsilon), \quad (\text{E.10})$$

for values of ε small enough that all cells remain positive. Here $\varepsilon = 0$ is the expansion point and $f(0) = I(P) = 0$. The second-order Taylor expansion is

$$f(\varepsilon) \approx \underbrace{f(0)}_{\text{zeroth order}} + \underbrace{f'(0)\varepsilon}_{\text{first order}} + \underbrace{\frac{f''(0)}{2!}\varepsilon^2}_{\text{second order}}. \quad (\text{E.11})$$

The first-order term is

$$f'(0) = \sum_{i,j} \{\ell_P(i,j) - 1\} h_{ij} = 0. \quad (\text{E.12})$$

Both parts vanish: $\ell_P(i,j) = 0$ at independence, and the probability constraint gives $\sum_{i,j} h_{ij} = 0$. Thus, positive and negative cell errors cannot create a linear change in MI around the independence table.

To state the surviving term, define the row and column changes $h_{i+} = \sum_j h_{ij}$ and $h_{+j} = \sum_i h_{ij}$. Differentiating the entropy form of MI twice gives

$$f''(0) = \sum_{i,j} \frac{h_{ij}^2}{p_{ij}} - \sum_i \frac{h_{i+}^2}{p_{i+}} - \sum_j \frac{h_{+j}^2}{p_{+j}}. \quad (\text{E.13})$$

Because $p_{ij} = p_{i+}p_{+j}$, define the part of each cell change that cannot be explained by changed margins as

$$a_{ij} = h_{ij} - p_{+j}h_{i+} - p_{i+}h_{+j}. \quad (\text{E.14})$$

Equation (E.13) can then be written as

$$f''(0) = \sum_{i,j} \frac{a_{ij}^2}{p_{i+}p_{+j}}. \quad (\text{E.15})$$

Substituting the zeroth-, first-, and second-order terms into Equation (E.11) gives

$$I(P_\varepsilon) = \underbrace{0}_{\text{zeroth and first order}} + \underbrace{\frac{\varepsilon^2}{2} \sum_{i,j} \frac{a_{ij}^2}{p_{i+}p_{+j}}}_{\text{second order}} + o(\varepsilon^2). \quad (\text{E.16})$$

The remainder $o(\varepsilon^2)$ is negligible relative to ε^2 as ε approaches zero.

For a multinomial sample, take $\varepsilon h = \hat{P} - P$. An empirical cell proportion has standard deviation $\sqrt{p_{ij}(1-p_{ij})/n}$, so each εh_{ij} is typically proportional to $n^{-1/2}$. The corresponding row- and column-adjusted errors εa_{ij} have the same scale. Equation (E.16) squares these errors, making $\hat{I}(P)$ proportional to n^{-1} under independence. Multiplying by $2n$ therefore produces a statistic with a non-vanishing scale:

$$2n\hat{I}(P) \xrightarrow{d} \chi_{(r-1)(c-1)}^2. \quad (\text{E.17})$$

This is the usual likelihood-ratio, or G , test of independence under a fixed alphabet and the standard positive-support large-sample conditions (Wilks, 1938; Brillinger, 2004). The chi-squared distribution describes the squared cell errors in

Equation (E.16). Expanded Welch instead standardises a signed first-order MI error on the $n^{-1/2}$ scale. At independence that signed term is absent, so the Student t distribution used by Expanded Welch is not the chi-squared distribution in another form.

Although squaring a regular asymptotic $N(0, 1)$ statistic gives a χ_1^2 limit, the present T is degenerate at independence. Its square is therefore not the likelihood-ratio G statistic.

The same likelihood-ratio connection appears for conditional information measures. For finite Markov chains, [Barnett and Bossomaier \(2012\)](#) relate transfer entropy to a log-likelihood ratio with an asymptotic chi-squared distribution; the wider information-theoretic setting is reviewed by [Bossomaier et al. \(2016\)](#). This is an adjacent application of the same second-order likelihood principle, not a two-sample equal-MI Welch test.

E.4 What alternative constructions produce

Using the observed margins to set $\hat{q}_{ij} = \hat{p}_{i+}\hat{p}_{+j}$ gives

$$2nD_{\text{KL}}(\hat{P}||\hat{Q}) = 2n\hat{I}(P), \quad (\text{E.18})$$

which is the likelihood-ratio statistic above. Because $\hat{Q}$ is constructed from the same sample as $\hat{P}$, it is not an independent second table and cannot be inserted into the two-sample Welch variance formula as though it supplied new information.

Repeatedly shuffling one variable preserves its observations while breaking the pairing with the other variable. The resulting empirical distribution is a permutation reference. Repeatedly drawing tables from $\hat{P}_X\hat{P}_Y$ is a parametric bootstrap reference. A single simulated Q table adds Monte Carlo variation but does not by itself define a calibrated null distribution.

The independent reference $Q = P_X P_Y$ is therefore mathematically valid, but it does not create a regular two-sample MI comparison. Analytic improvement of the sparse independence test would require a second-order finite-sample method and is not an immediate extension of the Expanded Welch calculation.

Bibliography

- Lionel Barnett and Terry Bossomaier. Transfer entropy as a log-likelihood ratio. *Physical Review Letters*, 109(13):138105, 2012. doi: 10.1103/PhysRevLett.109.138105.
- Robert M. Bell and Daniel F. McCaffrey. Bias reduction in standard errors for linear regression with multi-stage samples. *Survey Methodology*, 28(2):169–181, 2002. URL <https://www150.statcan.gc.ca/n1/pub/12-001-x/2002002/article/9058-eng.pdf>.
- Thomas B. Berrett and Richard J. Samworth. USP: An independence test that improves on Pearson’s chi-squared and the G -test. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 477(2256):20210549, 2021. doi: 10.1098/rspa.2021.0549.
- Terry Bossomaier, Lionel Barnett, Michael Harré, and Joseph T. Lizier. *An Introduction to Transfer Entropy: Information Flow in Complex Systems*. Springer, Cham, 2016. doi: 10.1007/978-3-319-43222-9.
- Richard P. Brent. *Algorithms for Minimization without Derivatives*. Prentice-Hall, Englewood Cliffs, New Jersey, 1973.
- David R. Brillinger. Some data analyses using mutual information. *Brazilian Journal of Probability and Statistics*, 18(2):163–183, 2004. URL <https://www.stat.berkeley.edu/~brill/Papers/MIBJPS.pdf>.
- EunYi Chung and Joseph P. Romano. Exact and asymptotically robust permutation tests. *Annals of Statistics*, 41(2):484–507, 2013. doi: 10.1214/13-AOS1090.
- W. Edwards Deming and Frederick F. Stephan. On a least squares adjustment of a sampled frequency table when the expected marginal totals are known. *Annals of Mathematical Statistics*, 11(4):427–444, 1940. doi: 10.1214/aoms/1177731829.
- Robert Drake and Apratim Guha. A mutual information-based k -sample test for

- discrete distributions. *Journal of Applied Statistics*, 41(9):2011–2027, 2014. doi: 10.1080/02664763.2014.899325.
- Bradley Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979. doi: 10.1214/aos/1176344552.
- Michael P. Fay and Barry I. Graubard. Small-sample adjustments for Wald-type tests using sandwich estimators. *Biometrics*, 57(4):1198–1206, 2001. doi: 10.1111/j.0006-341X.2001.01198.x.
- Apratim Guha and Tom Chothia. A two sample test based on mutual information. *Calcutta Statistical Association Bulletin*, 66(1–2):39–54, 2014. doi: 10.1177/0008068320140103.
- Frank R. Hampel. The influence curve and its role in robust estimation. *Journal of the American Statistical Association*, 69(346):383–393, 1974. doi: 10.1080/01621459.1974.10482962.
- Kermit Hutcheson. A test for comparing diversities based on the Shannon formula. *Journal of Theoretical Biology*, 29(1):151–154, 1970. doi: 10.1016/0022-5193(70)90124-4.
- Marcus Hutter. Distribution of mutual information. In *Advances in Neural Information Processing Systems 14*, pages 399–406, 2002. Conference held in 2001; arXiv:cs/0112019.
- Marcus Hutter and Marco Zaffalon. Distribution of mutual information from complete and incomplete data. *Computational Statistics & Data Analysis*, 48(3): 633–657, 2005. doi: 10.1016/j.csda.2004.03.010.
- Guido W. Imbens and Michal Kolesár. Robust standard errors in small samples: Some practical advice. *The Review of Economics and Statistics*, 98(4):701–712, 2016. doi: 10.1162/REST_a_00552.
- Arnold Janssen. Studentized permutation tests for non-i.i.d. hypotheses and the generalized Behrens–Fisher problem. *Statistics & Probability Letters*, 36(1):9–21, 1997. doi: 10.1016/S0167-7152(97)00043-6.
- Kirthivasan Kandasamy, Akshay Krishnamurthy, Barnabas Póczos, Larry Wasserman, and James M. Robins. Nonparametric von Mises estimators for entropies, divergences and mutual informations. In *Advances in Neural Information Processing Systems 28*, pages 397–405, 2015.

- Göran Kauermann and Raymond J. Carroll. A note on the efficiency of sandwich covariance matrix estimation. *Journal of the American Statistical Association*, 96(456):1387–1396, 2001. doi: 10.1198/016214501753382309.
- Marius Marinescu and Costel Balcau. On the use of mutual information for testing independence, 2025. arXiv:2502.17636.
- George A. Miller. Note on the bias of information estimates. In Henry Quastler, editor, *Information Theory in Psychology: Problems and Methods*, pages 95–100. Free Press, Glencoe, Illinois, 1955.
- Roelof Moddemeijer. On estimation of entropy and mutual information of continuous distributions. *Signal Processing*, 16(3):233–248, 1989. doi: 10.1016/0165-1684(89)90132-1.
- Roelof Moddemeijer. A statistic to estimate the variance of the histogram-based mutual information estimator based on dependent pairs of observations. *Signal Processing*, 75(1):51–63, 1999. doi: 10.1016/S0165-1684(98)00224-2.
- Ricardo Mora and Javier Ruiz-Castillo. The statistical properties of the mutual information index of multigroup segregation. Working Paper, Economics Series 48 09-84, Universidad Carlos III de Madrid, Departamento de Economía, 2009.
- Ricardo Mora and Javier Ruiz-Castillo. Entropy-based segregation indices. *Sociological Methodology*, 41(1):159–194, 2011. doi: 10.1111/j.1467-9531.2011.01237.x.
- Wei Pan and Melanie M. Wall. Small-sample adjustments in using the sandwich variance estimator in generalized estimating equations. *Statistics in Medicine*, 21(10):1429–1441, 2002. doi: 10.1002/sim.1142.
- Liam Paninski. Estimation of entropy and mutual information. *Neural Computation*, 15(6):1191–1253, 2003. doi: 10.1162/089976603321780272.
- Mark S. Roulston. Estimating the errors on measured entropy and mutual information. *Physica D: Nonlinear Phenomena*, 125(3–4):285–294, 1999. doi: 10.1016/S0167-2789(98)00269-3.
- Franklin E. Satterthwaite. An approximate distribution of estimates of variance components. *Biometrics Bulletin*, 2(6):110–114, 1946. doi: 10.2307/3002019.
- Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27:379–423, 623–656, 1948. Parts I and II.

- Aad W. van der Vaart. *Asymptotic Statistics*. Cambridge University Press, Cambridge, 1998. doi: 10.1017/CBO9780511802256.
- Abraham Wald. Tests of statistical hypotheses concerning several parameters when the number of observations is large. *Transactions of the American Mathematical Society*, 54(3):426–482, 1943. doi: 10.1090/S0002-9947-1943-0012401-3.
- Bernard L. Welch. The generalization of “Student’s” problem when several different population variances are involved. *Biometrika*, 34(1–2):28–35, 1947. doi: 10.1093/biomet/34.1-2.28.
- Samuel S. Wilks. The large-sample distribution of the likelihood ratio for testing composite hypotheses. *The Annals of Mathematical Statistics*, 9(1):60–62, 1938. doi: 10.1214/aoms/1177732360.
- Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. doi: 10.1080/01621459.1927.10502953.