

Finite-Sample Differential Mutual Information Inference: A Bias-Corrected Welch-Type Test and Variance Audit

Daniel Zhao

Master's Research in Information Theory

Draft methods article

3 August 2026

Abstract

Mutual information (MI) quantifies arbitrary statistical dependence, yet standard significance procedures are primarily designed to test independence within one population. A different scientific question arises when two independent populations may have distinct joint distributions but the same strength of dependence: $H_0 : I(P) = I(Q)$, allowing $P \neq Q$. Resampling can address this weak null but is computationally costly, whereas a bias-corrected influence-function Wald test is fast but can be liberal in finite samples. We study a deterministic bias-corrected Welch-type differential-MI test (BCW-DMI) that combines the classical first-order correction for discrete plug-in MI, an empirical influence-function standard error, and a Student reference. This construction is an MI extension of the Hutcheson entropy-comparison template, not a wholly new statistical form.

The method was evaluated on 1.22 million weak-null table pairs and 50,000 power pairs spanning square and rectangular alphabets from 2×2 to 20×20 , balanced and skewed margins, and sample-size ratios up to 1:4. At nominal $\alpha = 0.05$, mean absolute false-positive-rate error fell from 0.01177 to 0.01084 in the targeted hard grid, a 7.9% relative improvement, while broad-grid error changed from 0.00514 to 0.00504. An independently generated holdout confirmed the direction of the effect. Mean power loss was 0.00154, and median runtime was 0.128 ms versus 0.117 ms for the normal test. Studentized permutation remained better calibrated (hard-grid error 0.00642) but was substantially slower. An adversarial audit showed that the implemented component degrees of freedom are heuristic for plug-in influence variances. A post-hoc 960,000-pair diagnostic found that a full variance-functional influence calculation reduced hard-grid error to 0.00823 and reproduced the direction on six newly generated hard populations (0.00993 to 0.00662), but was not uniformly better across 72 regular scenarios. The evidence supports a low-cost research direction, not an exact Welch theorem or a universal sparse-table solution.

Keywords: mutual information; differential mutual information; Welch-Satterthwaite approximation; influence function; finite-sample inference; permutation

test

1 Introduction

Mutual information (MI) is a model-free measure of statistical dependence introduced by Shannon [1]. For categorical variables it summarizes the departure of a joint distribution from the product of its margins and detects associations that need not be linear, monotone, or low-dimensional. These properties make MI attractive in neuroscience, genomics, ecology, remote sensing, and complex-systems research [2, 3].

Most familiar MI significance procedures answer a one-population question: are X and Y independent, or equivalently, is $I(X; Y) = 0$? The likelihood-ratio statistic $2n\hat{I}$ has an asymptotic chi-squared reference for fixed categorical alphabets, and permutation procedures commonly shuffle one variable to destroy dependence. Toolkits such as JIDT make these calculations accessible [3].

Many comparative studies require a different test. Suppose the same categorical variables are measured in two independent populations, such as treatment and control, two environments, or two time periods. The question is then

$$H_0 : I(P) = I(Q) \quad \text{against} \quad H_1 : I(P) \neq I(Q), \quad (1)$$

where the full joint distributions P and Q may differ. This is a *weak null*: equal MI does not imply equal margins, cell probabilities, or association patterns. A chi-squared independence test in either table does not answer Eq. (1). Likewise, JIDT's standard shuffled independence significance API is not a direct comparator.

The usual analytic route treats MI as a smooth multinomial functional away from independence. A delta-method or influence-function calculation gives an asymptotically normal estimator, and two independent variance contributions can be added. For finite samples, however, the plug-in MI estimate is biased upward and its estimated uncertainty is itself noisy. Unequal sample sizes make the leading bias differ between groups. A normal-reference Wald test can therefore reject too often in skewed or low-expected-count regimes.

This article studies a deliberately small modification, denoted BCW–DMI. It first applies the classical $O(n^{-1})$ bias correction to each MI estimate, studentizes their difference using empirical local-information variances, and then replaces the standard-normal reference by a Student distribution whose effective degrees of freedom use the Welch–Satterthwaite formula [4, 5]. The conceptual pipeline is shown in Fig. 1. Because every operation is computed directly from two count tables, the method is deterministic and has $O(rc)$ cost once the tables are formed.

The contribution is a transport, audit, and evaluation rather than the invention of a test template. Hutcheson’s entropy test already combines a difference of Shannon functionals, delta-method variances, and a Satterthwaite reference [6]. The present work asks what happens after transporting that architecture to discrete MI, adding its specific leading bias correction, and evaluating it under the unrestricted weak null in Eq. (1). The degrees-of-freedom audit then identifies why the textbook substitution is incomplete and motivates an MI-specific variance-functional refinement.

2 Related Work

2.1 Finite-sample estimation of mutual information

For an $r \times c$ discrete distribution $P = (p_{ij})$, MI is

$$I(P) = \sum_{i=1}^r \sum_{j=1}^c p_{ij} \log \frac{p_{ij}}{p_{i+} p_{+j}}. \quad (2)$$

Replacing probabilities by empirical proportions gives the maximum-likelihood or plug-in estimator. Its simplicity hides a systematic finite-sample upward bias. Miller’s entropy correction [7], subsequent MI-specific calculations [8], and broader analyses [9] show why sparse sampling can contaminate information estimates even after a leading correction.

For fixed positive support, the plug-in MI estimator is asymptotically normal away from independence. Its first-order variance is the variance of the local log density ratio [8]. Influence-function methods provide a general language for this expansion and for information-theoretic functionals more broadly [10]. These results motivate fast Wald inference, but the first-order influence function degenerates at exact independence, where the chi-squared rather than normal limit is relevant.

2.2 Comparing mutual information values

Hutcheson’s classical test compares two Shannon diversity values by dividing their difference by estimated entropy variances and using Satterthwaite degrees of freedom [6]. Its modern implementation uses the same

broad structure as Eqs. (9) and (12), with n rather than $n - 1$ in the component denominators. Extending it to MI is not automatic: MI depends jointly on cell and marginal probabilities, has leading bias $d/(2n)$, and has a degenerate first-order influence function at independence. Nevertheless, Hutcheson is the direct template and defines the appropriate novelty boundary.

The idea of comparing MI across populations is also established. Mora and Ruiz-Castillo derived asymptotic inference and pairwise comparisons for the discrete MI segregation index in a working-paper line [11]; their published account establishes the wider role and properties of the M index [12]. Related segregation research shows how bias adjustments, asymptotic tests, and bootstrap refinements can be combined for nonlinear indices [13]. Hart and Giszter combined jackknife uncertainties when comparing histogram MI estimates in neuroscience [14].

Other studies apply familiar mean tests to *replicated* MI estimates. Sarkar and Pandey used a Student test on MI values obtained from randomized or shuffled astronomical data [15]. Prince et al. applied Welch comparisons to per-neuron MI outcomes [16]. Martin et al. used an unequal-variance Welch test on repeated, non-overlapping partition estimates from a continuous KSG estimator [17]; the estimator itself builds on controlled-bias KSG work such as Holmes and Nemenman [18]. These studies establish that Welch testing has entered MI applications. They do not, however, derive the two variance components from one discrete multinomial table per population as done here. Mather et al. evaluated Welch and MI leakage detectors side by side, but did not apply the Welch reference to the MI estimator itself [19].

Resampling is an essential comparator rather than an afterthought. Bootstrap methods can improve nonlinear segregation-index inference [13], and simulation studies of Shannon diversity report strong coverage for bootstrap- t intervals [20]. A multinomial bootstrap- t comparison for differential MI is therefore required in the next confirmatory study.

2.3 Permutation under a weak null

Raw permutation is exact when observations are exchangeable under the strong null $P = Q$. Under Eq. (1), P and Q may differ in every respect other than MI, so group labels are generally not exchangeable. Consequently, naively permuting labels can test the wrong null. General permutation theory shows that studentization can recover asymptotic validity for equality of parameters under weak assumptions while retaining finite exactness when the complete distributions are equal [21]. We therefore use a studentized table-label permutation as the resampling benchmark, rather than raw permutation or a one-table independence shuffle.



$$H_0 : I(P) = I(Q) \text{ while allowing } P \neq Q$$

Figure 1: **Overview of BCW-DMI.** Two independent contingency tables are converted to plug-in MI estimates, corrected for leading finite-sample bias, studentized with influence-function variances, and compared with a Welch-type Student reference. The target is equality of MI, not equality of the complete distributions and not independence within either table.

3 Bias-Corrected Welch-Type Test

3.1 Data and estimand

Let $N^P = (N_{ij}^P)$ and $N^Q = (N_{ij}^Q)$ be independent $r \times c$ count tables with totals n_P and n_Q . Categories must be aligned between groups. Write $\hat{p}_{ij} = N_{ij}^P / n_P$ and define $\hat{q}_{ij}$ analogously. All formulas use natural logarithms and therefore report nats. Changing to bits multiplies the estimated difference and its standard error by the same constant, leaving the test statistic and p-value unchanged.

The parameter of interest is the signed differential MI

$$\Delta = I(P) - I(Q). \quad (3)$$

The test is two-sided by default, although a directional alternative follows immediately when it is scientifically pre-specified.

3.2 Bias-corrected effect estimate

Let $d = (r-1)(c-1)$. Under fixed positive support, the leading plug-in MI bias is $d/(2n)$. We therefore define

$$\tilde{I}_P = \hat{I}_P - \frac{d}{2n_P}, \quad (4)$$

$$\tilde{I}_Q = \hat{I}_Q - \frac{d}{2n_Q}, \quad (5)$$

$$\hat{\Delta}_{BC} = \tilde{I}_P - \tilde{I}_Q. \quad (6)$$

This correction is especially important when $n_P \neq n_Q$, because the uncorrected difference has leading bias $d(1/n_P - 1/n_Q)/2$ even under equal population MI.

The correction uses the declared alphabet dimensions, not the observed nonempty dimensions. It is appropriate for the stated fixed-positive-support asymptotics. In a severely sparse sample it can overcorrect, and corrected MI values may be negative. Neither clipping nor data-dependent support changes are applied

because both would alter the estimand and sampling expansion.

3.3 Influence-function uncertainty

For each positive cell, define local information

$$\ell_{ij}^P = \log \frac{\hat{p}_{ij}}{\hat{p}_{i+}\hat{p}_{+j}}. \quad (7)$$

The empirical influence variance is

$$\hat{V}_P = \sum_{ij: \hat{p}_{ij} > 0} \hat{p}_{ij} \left(\ell_{ij}^P - \hat{I}_P \right)^2, \quad (8)$$

with $\hat{V}_Q$ defined analogously. Independence of the samples gives

$$\widehat{SE} = \sqrt{\frac{\hat{V}_P}{n_P} + \frac{\hat{V}_Q}{n_Q}}. \quad (9)$$

The standardized statistic is

$$T = \frac{\hat{\Delta}_{BC}}{\widehat{SE}}. \quad (10)$$

Equations (6)–(10) are identical to the bias-corrected normal Wald baseline.

3.4 Welch-Satterthwaite reference

Set

$$a = \frac{\hat{V}_P}{n_P}, \quad b = \frac{\hat{V}_Q}{n_Q}. \quad (11)$$

BCW-DMI assigns the effective degrees of freedom

$$\hat{\nu} = \frac{(a+b)^2}{a^2/(n_P-1) + b^2/(n_Q-1)}. \quad (12)$$

The two-sided p-value and $100(1-\alpha)\%$ confidence interval are

$$p_{BCW} = 2 \Pr\{t_{\hat{\nu}} \geq |T|\}, \quad (13)$$

$$\hat{\Delta}_{BC} \pm t_{\hat{\nu}, 1-\alpha/2} \widehat{SE}. \quad (14)$$

Table 1: **Operational algorithm for one comparison.**

1	Validate two aligned, independent count tables.
2	Compute plug-in MI in each table using Eq. (2).
3	Subtract $d/(2n)$ and form the difference in Eq. (6).
4	Compute local-information variances using Eq. (8).
5	Form the standard error and statistic in Eqs. (9)–(10).
6	Compute $\hat{\nu}$ from Eq. (12).
7	Return the Student p-value, confidence interval, and support diagnostics.

For finite $\hat{\nu}$, Student tails are heavier than normal tails. Thus, at a fixed T , p_{BCW} is never smaller than the normal Wald p-value. The method is a transparent conservative adjustment, not a new effect estimator.

3.5 First-order justification and qualification

For a positive fixed table,

$$\sqrt{n}\{\hat{\mathbf{I}} - \mathbf{I}(P)\} = \frac{1}{\sqrt{n}} \sum_{k=1}^n \psi_P(X_k, Y_k) + o_P(1), \quad (15)$$

where $\psi_P(x, y) = \log\{p_{xy}/(p_{x+}p_{+y})\} - \mathbf{I}(P)$ is the MI influence function. Combining the independent expansions for P and Q establishes the normal-reference Wald limit.

Under fixed finite alphabets, positive population cell probabilities, independent and identically distributed observations within each group, independent groups, sample-size proportions bounded away from zero and one, and positive influence variances, Eq. (15) implies that T converges in distribution to $N(0, 1)$ under H_0 . Moreover, $\hat{\nu} \rightarrow \infty$, so the Student reference converges to the same normal limit. This is asymptotic compatibility by construction, not a second-order accuracy result.

Classical Welch inference combines ordinary sample variances with known Gaussian degrees-of-freedom interpretations. Here, $\hat{V}_P$ and $\hat{V}_Q$ are nonlinear plug-in functionals: the local scores change with the same empirical tables used to estimate MI. Assigning $n_P - 1$ and $n_Q - 1$ in Eq. (12) is therefore a Welch–Satterthwaite *analogy*. Its value must be evaluated empirically.

3.6 Worked numerical example

Consider the following synthetic 3×3 tables:

$$N^P = \begin{pmatrix} 45 & 15 & 10 \\ 15 & 45 & 10 \\ 10 & 10 & 40 \end{pmatrix}, \quad N^Q = \begin{pmatrix} 48 & 32 & 20 \\ 32 & 48 & 20 \\ 20 & 20 & 40 \end{pmatrix}.$$

Their totals are $n_P = 200$ and $n_Q = 280$, and $d = 4$. The plug-in estimates are 0.211316 and 0.054303 nats. Subtracting 0.010000 and 0.007143 gives corrected estimates 0.201316 and 0.047160, hence $\hat{\Delta}_{\text{BC}} = 0.154155$ nats.

The empirical influence variances are $\hat{V}_P = 0.395412$ and $\hat{V}_Q = 0.110863$. Therefore,

$$\widehat{\text{SE}} = \sqrt{0.395412/200 + 0.110863/280} = 0.048713,$$

and $T = 3.16454$. Equation (12) gives $\hat{\nu} = 278.71$. The normal p-value is 0.001553, whereas BCW–DMI returns $p = 0.001725$ and a 95% interval $[0.05826, 0.25005]$ nats. Both references reject equality, but the Student reference is slightly more conservative. The minimum independence-expected counts are 18.0 and 22.9, so this example lies comfortably inside the dense regular regime. It illustrates the calculation only and is not part of the simulation evidence.

3.7 Complexity and recommended output

Given two count tables, plug-in MI, influence variance, and support diagnostics each require a constant number of scans over rc cells. Consequently, time is $O(rc)$ and additional working memory is $O(rc)$ in the vectorized implementation. If raw observation pairs are supplied, forming the tables first adds $O(n_P + n_Q)$ time. No random number generator, resample count, or convergence tolerance is required.

A complete analysis should report n_P, n_Q, r, c , the two plug-in and bias-corrected MI estimates, $\hat{\Delta}_{\text{BC}}$, $\widehat{\text{SE}}$, T , $\hat{\nu}$, the Student p-value and confidence interval, and the normal-reference p-value as a comparator. Zero-cell fractions and expected-count diagnostics should accompany sparse applications. A noncomputable first-order variance must be reported as invalid, not silently converted to $p = 1$. The bias correction and the Welch reference are distinct: the former changes the point estimate, while the latter changes only the inferential reference.

4 Experimental Design

4.1 Weak-null population construction

Every null experiment used independent multinomial samples from population pairs satisfying $\mathbf{I}(P) = \mathbf{I}(Q)$ to numerical tolerance while allowing different margins and joint distributions. Random Dirichlet margins and interaction structures generated heterogeneous tables, and association strength was adjusted to target MI levels of 0.03, 0.07, or 0.15 nats. Common sampled table pairs were supplied to all methods, producing paired method comparisons without training, threshold fitting, or cross-method data leakage.

4.2 Decisive simulation

The main experiment contained four stages:

1. A broad grid of 144 population pairs, shapes from 2×2 to 20×20 , square and rectangular tables, equal and unequal samples with ratios 1:1, 1:2, and 1:4, and 5,000 replicates per pair.
2. A targeted hard grid of 12 low-density, unequal-sample population pairs across six shapes, with 20,000 fresh replicates per pair.
3. A 26-scenario stress grid with sample sizes 20–200 and 10,000 replicates per scenario. This deliberately approached or crossed the first-order boundary.
4. Five 3×3 power scenarios with true differences 0.02, 0.05, and 0.10 nats and 10,000 replicates each.

In total, the decisive run evaluated 1,220,000 null table pairs and 50,000 power pairs. Twelve thousand hard-grid pairs also received 999 optimized studentized table permutations.

4.3 Independent adversarial holdout

After the decisive results were inspected, a separate audit fixed new population, simulation, and bootstrap seeds. It generated 72 fresh weak-null population pairs with 10,000 table pairs each, re-evaluated six frozen hard designs with 10,000 fresh table pairs each, and generated 72 strong-null $P = Q$ scenarios with 5,000 table pairs each. This holdout added 1.14 million table-pair evaluations and was designed to test the direction and generalizability of the observed effect rather than tune the method.

4.4 Post-hoc variance diagnostic

A subsequent methodological audit fixed a third population seed before simulation. It generated 20,000 new table pairs for each of the 12 decisive hard populations and 10,000 pairs for each of 72 entirely new weak-null populations, totalling 960,000 comparisons. It measured $E(\hat{V})/V$, the empirical variance of $\hat{\Delta}_{BC}$, and its correlation with the estimated variance. It also compared Eq. (12) with local-score-kurtosis degrees of freedom, full variance-functional influence degrees of freedom, and a diagnostic normal test using the known population variance. Because this audit was designed after inspecting the original study, it is explicitly post-hoc even though its new population and simulation seeds were fixed before execution.

4.5 Comparators and metrics

The primary analytic comparator was the bias-corrected normal Wald test, which shares the estimate, standard error, and statistic with BCW–DMI. The resampling comparator was a studentized analytic table-label permutation test. An exploratory variant using $n - 1$ rather than n in the empirical variance

denominator was reported separately but was not used to define the primary method.

At nominal $\alpha = 0.05$ and 0.10, we recorded false-positive rate (FPR), Wilson intervals, mean absolute FPR error (MAE), confidence-interval coverage, valid-result rate, paired rejection changes, power, and runtime. Calibration is the central criterion: at $\alpha = 0.05$, an ideal method has FPR 0.05, and smaller $|\text{FPR} - 0.05|$ is better. Candidate methods were applied to the same table pair in every replicate. Consequently, their rejection changes are paired; the larger marginal Monte Carlo uncertainty of either FPR does not describe the uncertainty of the much smaller within-replicate difference.

5 Results

5.1 Calibration

Table 2 summarizes the decisive null experiment. In the broad regular grid, both analytic methods were well calibrated: mean FPR was 0.04977 for normal Wald and 0.04951 for BCW–DMI. MAE improved slightly from 0.00514 to 0.00504, and the proportion of scenarios within the practical 0.035–0.065 band increased from 97.22% to 97.92%.

The targeted hard grid was deliberately liberal under normal Wald. Its mean FPR fell from 0.06177 to 0.06084, and MAE fell from 0.01177 to 0.01084, a 7.9% relative reduction. Welch removed 224 normal-Wald rejections among 240,000 hard-null pairs; no Welch-only rejection is possible because the Student p-value is nested above the normal p-value. Coverage increased by 0.00093.

The stress-grid result is negative at the scenario level. Although mean FPR moved toward 0.05, the proportion inside the practical calibration band fell from 19.23% to 15.38%. The correction is monotone conservative, so already conservative 2×2 settings became worse; the present method has no validated diagnostic that identifies such cases in advance. In addition, 34 first-order-degenerate cases were explicitly returned as invalid. The 99.987% valid rate is high, but stress-grid FPR is conditional on a computable positive standard error.

5.2 Independent holdout

The holdout reproduced the direction of the effect without reusing the broad population pairs (Table 3). In the fresh weak-null broad grid, mean MAE gain was 0.000210 with a scenario-bootstrap 95% interval [0.000090, 0.000344]. In all six frozen hard designs, BCW–DMI improved calibration; mean gain was 0.001283 [0.000917, 0.001650]. The fresh strong-null grid also showed a small positive average gain. These results support a real conservative effect, but not uniform superiority: in the fresh broad grid, 26 scenarios improved, 14 worsened, and 32 tied.

Table 2: **Decisive null calibration at nominal $\alpha = 0.05$.** The stress grid is outside or near the intended first-order operating boundary. MAE is mean absolute false-positive-rate error.

Grid	Method	Scenarios	Mean FPR	MAE	95% coverage	In 3.5–6.5% band
Broad	Normal Wald	144	0.04977	0.00514	0.95023	97.22%
Broad	BCW–DMI	144	0.04951	0.00504	0.95049	97.92%
Hard	Normal Wald	12	0.06177	0.01177	0.93823	75.00%
Hard	BCW–DMI	12	0.06084	0.01084	0.93916	91.67%
Stress	Normal Wald	26	0.06580	0.03180	0.93420	19.23%
Stress	BCW–DMI	26	0.06276	0.03037	0.93724	15.38%




Figure 2: **Paired calibration in the decisive simulation.** Left: hard-grid FPRs for the normal baseline, primary Welch reference, and exploratory unbiased-variance sensitivity. The dashed line is nominal 0.05 and the shaded band is 0.035–0.065. Right: each point is one broad scenario; points below the diagonal indicate a lower rejection rate under BCW–DMI. The adjustment is intentionally small because hard-grid median effective degrees of freedom ranged from approximately 186 to 812.

5.3 Permutation comparison

Studentized permutation remained the most accurate comparator in the hard anchors (Table 4). Its mean FPR was 0.05458 and MAE was 0.00642, versus 0.06083 and 0.01100 for BCW–DMI on the same 12 anchor sets. Every permutation anchor was inside the practical calibration band. Thus, BCW–DMI closes only part of the gap between normal analytic inference and resampling; it does not replace permutation when maximum finite-sample calibration is worth the additional computation.

5.4 Power and computation

The heavier Student tail produces a predictable but small power cost. Across five power scenarios, mean absolute loss relative to normal Wald was 0.00154; the largest scenario-level loss was 0.0032 at $n_P = n_Q = 150$. For a 0.05-nat effect at $n = 300$, power changed from 0.2776 to 0.2761. For a 0.10-nat effect, it changed from 0.7438 to 0.7419.

Median single-pair runtime was 0.117 ms for normal Wald and 0.128 ms for BCW–DMI, a 9.1% relative

overhead but only 0.011 ms in absolute terms. Across the 12 permutation anchors, median time for 999 permutations was 3.006 ms, with a range from roughly 1.1 to 5.9 ms depending on table shape. The implementation is vectorized, so the exact ratio depends on batching and hardware. The key distinction is algorithmic: the analytic calculation scans the rc cells once, whereas permutation repeats table allocation and studentized MI evaluation B times.

6 Adversarial Audit

6.1 Correctness and reproducibility

The implementation passed direct hand calculations of MI, bias, variance, degrees of freedom, statistic, and p-value. Swapping groups and relabelling rows or columns left the p-value unchanged. Student p-values were never smaller than their normal counterparts, and the two references converged at large sample size. Natural-log units were used consistently. All valid p-values were finite and inside $[0, 1]$; malformed tables and degenerate first-order variances were rejected

Table 3: **Adversarial holdout at nominal $\alpha = 0.05$.** Intervals are scenario-bootstrap intervals for the mean MAE gain ($\text{MAE}_N - \text{MAE}_{\text{BCW}}$).

Holdout stage	Populations	Replicates	Normal MAE	BCW MAE	Mean gain [95% interval]
Fresh weak-null broad	72	10,000	0.004792	0.004582	0.000210 [0.000090, 0.000344]
Frozen hard design	6	10,000	0.009083	0.007800	0.001283 [0.000917, 0.001650]
Fresh strong null	72	5,000	0.005142	0.004958	0.000183 [0.000064, 0.000322]




Figure 3: **Accuracy and cost summary.** Left: normal and BCW-DMI MAE across the decisive study and independent holdout. The largest benefit occurs in deliberately liberal hard regimes; broad and strong-null effects are small. Right: median observed single-pair runtime on a logarithmic scale. The permutation result uses 999 optimized table permutations and is not JIDT’s one-table independence shuffle.

Table 4: **Hard-grid permutation anchors.** Each of 12 scenarios used 1,000 table pairs and 999 studentized permutations.

Method	Mean FPR	MAE	In band
Normal Wald	0.06208	0.01208	66.67%
BCW-DMI	0.06083	0.01100	75.00%
Studentized permutation	0.05458	0.00642	100.00%

explicitly.

Simulation streams were separated across population generation, table sampling, permutation, and power experiments. The weak null held to a maximum saved numerical discrepancy of 9.3×10^{-14} nats. Reconstructing saved replicates from scenario and simulation seeds reproduced estimates and p-values to floating-point precision. No fitted hyperparameter, learned threshold, or training stage consumed validation outcomes.

6.2 Degrees-of-freedom limitation

The central theoretical qualification is visible in Fig. 4. In five representative populations, the implemented naive total degrees of freedom ranged from 171 to

858, whereas empirical moment-matched values ranged from 28 to 196. A separately derived influence function for the variance functional predicted 26 to 181 and tracked the empirical values more closely.

This discrepancy does not invalidate the first-order asymptotic result: every candidate reference converges to normal as sample size grows. It does show that Eq. (12) is not literally moment matching the sampling distribution of the nonlinear influence-variance estimator. Correlation between the estimated MI difference and its total estimated variance reached 0.84 in a skewed 2×2 audit, so improving the denominator degrees of freedom alone may still not yield an exact Student law.

6.3 Variance bias and refined degrees of freedom

The post-hoc diagnostic rejected a simple universal variance-bias explanation. Mean estimated SE^2 divided by its population first-order value was 1.001 on the decisive hard grid and 1.066 on the fresh 72-scenario grid, while the empirical variance of the corrected MI difference was 1.071 and 1.063 times that population value. Even the oracle-variance normal test was not uniformly calibrated. Thus, finite-sample




Figure 4: **The implemented degrees of freedom are heuristic.** Empirical moment degrees of freedom for the total estimated variance are well below Eq. (12). A variance-functional influence calculation tracks the audit more closely. The current procedure must therefore be described as Welch-inspired, not exact.

numerator variation and numerator–denominator dependence matter alongside variance estimation.

For completeness, let $V(P) = \text{Var}_P(\ell)$, $m_2 = E_P(\ell^2)$, and $\mu = I(P)$. The influence function of the variance functional is

$$g_P(x, y) = \ell_{xy}^2 - m_2 + 2\{\ell_{xy} - E(\ell | X = x) - E(\ell | Y = y) + \mu\} - 2\mu(\ell_{xy} - \mu). \quad (16)$$

Writing $\tau_P^2 = \text{Var}_P\{g_P(X, Y)\}$, first-order moment matching gives $\nu_{V,P} = 2n_P V(P)^2 / \tau_P^2$. Plug-in versions for P and Q can replace $n_P - 1$ and $n_Q - 1$ in Eq. (12) at unchanged $O(rc)$ order.

Table 5 shows why this candidate deserves a new confirmatory study. It materially improved the deliberately liberal hard regimes and had its largest gain at $\alpha = 0.01$. Across all fresh regular scenarios, however, it improved 28, worsened 39, and tied five at $\alpha = 0.05$; the simpler local-kurtosis approximation had slightly lower mean MAE there. All candidate valid-result rates were one in this audit.

6.4 Post-protocol decision

The frozen validation protocol used an internal rule requiring at least a 20% reduction in hard-grid MAE. The observed 7.9% reduction did not satisfy that rule, and the generated metadata correctly records NO-GO. The variance-functional result was obtained later and cannot retroactively change that decision. It instead defines a new candidate requiring a separately frozen protocol. This distinction preserves the original numerical result, comparators, and report while allowing the research direction to evolve.

7 Discussion

7.1 What the method adds

BCW–DMI establishes a transparent MI-specific baseline between normal analytic inference and studentized permutation. It adds almost no computational burden and provides a small conservative adjustment where the normal reference is mildly liberal, but Hutcheson’s entropy test makes clear that this architecture is transported rather than invented. The stronger methodological contribution is the variance audit: it explains why assigning $n - 1$ to a nonlinear plug-in variance is unjustified and supplies a computable $O(rc)$ alternative through Eq. (16).

The size of the improvement is mathematically unsurprising. In the hard grid, median effective degrees of freedom were approximately 186–812. Student critical values at those degrees of freedom are close to normal critical values, so few decisions change. The variance-functional candidate lowers those degrees of freedom in the regimes where estimated uncertainty is most variable. Its nonuniform broad-grid behaviour shows that this is still an approximation to a dependent studentized statistic, not a universal pivot.

7.2 Relationship to standard MI tests

Table 6 separates four procedures that are easily conflated. Chi-squared and JIDT shuffling address independence within one population. They are appropriate for $H_0 : I(P) = 0$, not Eq. (1). Normal Wald, BCW–DMI, and studentized two-sample permutation address equal MI across independent populations. Raw group-label permutation is exact only under the stronger $P = Q$ null; studentization is required for the heterogeneous weak null.

7.3 Novelty boundary

The original construction claim must be narrowed. Hutcheson introduced the entropy analogue in 1970 [6]; Mora and Ruiz-Castillo developed asymptotic comparison for the MI segregation index [11]; and published examples use Welch on repeated continuous-MI estimates [17] and per-unit MI outcomes [16]. Nor are the plug-in estimator, $d/(2n)$ correction, influence variance, studentization, or Satterthwaite formula individually new.

The defensible contribution is therefore methodological transport plus audit: an explicit Hutcheson-type procedure for two independent discrete MI values under the unrestricted weak null, unusually broad paired finite-sample evaluation, and the derivation and preliminary test of variance-functional degrees of freedom tailored to the MI uncertainty estimator. The last element is more distinctive than the $n - 1$ baseline but remains post-hoc. Before any publication-priority statement, the search should be repeated in authenticated Scopus, Web of Science, MathSciNet, zbMATH,

Table 5: **Post-hoc degrees-of-freedom diagnostic.** Entries are mean absolute FPR errors. Variance IF uses Eq. (16); the oracle uses known population variances and is diagnostic, not usable.

Population set (α)	Normal	$n - 1$ Welch	Local kurtosis	Variance IF	Oracle
Decisive hard (0.05)	0.01270	0.01169	0.01040	0.00823	0.00785
Decisive hard (0.01)	0.00568	0.00506	0.00428	0.00266	0.00354
Fresh hard (0.05)	0.00993	0.00912	0.00775	0.00662	0.00837
Fresh hard (0.01)	0.00407	0.00360	0.00317	0.00232	0.00317
Fresh all (0.05)	0.00463	0.00445	0.00423	0.00446	0.00756

Table 6: **How the proposed procedure relates to common significance methods.** A method is only a valid comparator when it targets the same null.

Method	Null hypothesis	Computation	Role
Chi-squared / JIDT analytic	$I(P) = 0$	Closed-form reference	One-table independence; not a test of equal MI across groups.
JIDT variable shuffling	$I(P) = 0$	Repeated data-label shuffles	One-table empirical independence significance; not a direct comparator.
Bias-corrected normal Wald	$I(P) = I(Q)$	$O(rc)$, deterministic	Historical two-sample analytic baseline.
BCW-DMI	$I(P) = I(Q)$	$O(rc)$, deterministic	Hutcheson-type MI baseline with heuristic $n - 1$ component df.
Variance-IF candidate	$I(P) = I(Q)$	$O(rc)$, deterministic	Post-hoc MI-specific moment refinement requiring confirmation.
Studentized group permutation	$I(P) = I(Q)$ asymptotically; exact if $P = Q$	B repeated allocations and evaluations	More accurate regular resampling benchmark, with higher cost.

and ProQuest with Hutcheson and diversity-index terms included.

7.4 Limitations and operating range

The method requires independent observations within each group and independent groups; fixed, aligned categorical alphabets; positive population support; sample-size proportions bounded away from zero and one; and nonzero population influence variance. It does not currently cover paired, clustered, longitudinal, or time-series samples.

Exact and near independence are outside scope because the first-order MI influence function degenerates. Structural zeros, unknown support, severe sampling sparsity, or alphabets growing with sample size can invalidate both the leading bias correction and variance approximation. The reported zero-cell and expected-count diagnostics describe a table but are not proven validity rules. Conditional MI, transfer entropy, and continuous MI require new derivations.

The simulation evidence is extensive in null calibration but less broad in power: the five alternatives are 3×3 tables with equal sample sizes. The permutation comparison covers 12 hard anchors with 1,000 outer replicates each. A final thesis study should broaden power across shapes and unequal samples, increase the outer permutation replicates, pre-specify $\alpha = 0.01$, add multinomial bootstrap- t , test declared versus observed support dimensions, and include real-data demonstrations. These additions improve ex-

ternal validity; simply increasing the existing null replicate counts would add little.

7.5 Next methodological step

The next step is a pre-specified head-to-head study of the local-kurtosis and full variance-functional references. Equation (16) tracked empirical component degrees of freedom more closely than $n - 1$ and produced the stronger hard-grid result, whereas local kurtosis had slightly better average fresh-grid MAE. Neither may replace Eq. (12) until evaluated on another untouched population family against studentized permutation and bootstrap- t . Corrected Satterthwaite formulas proposed for general variance synthesis may also be relevant [22], but their assumptions are not yet established for MI.

8 Conclusion

BCW-DMI transports the Hutcheson-Satterthwaite comparison architecture to a distinct information-theoretic problem: equality of two independent discrete MI values under a weak null. Leading plug-in bias correction and MI-specific influence uncertainty make the transport substantive, but the $n - 1$ reference itself is only mildly beneficial. Studentized permutation remained more accurate.

The audit changes the research emphasis. The estimated MI difference and its plug-in uncertainty are

dependent, the ordinary component degrees of freedom are too large, and no single variance rescaling repairs calibration. A full variance-functional influence calculation is deterministic, remains $O(rc)$, and produced materially lower hard-regime error in a new diagnostic, especially at $\alpha = 0.01$. Its nonuniform regular-grid behaviour demands a fresh confirmatory study. The defensible thesis is thus not a new Welch test, but an MI-specific transport, failure analysis, and promising moment refinement with clearly defined limits.

Reproducibility and Data Availability

The implementation, frozen protocol, experiment runners, aggregate result tables, seeds, environment metadata, tests, and figure-generation code are stored in the accompanying research repository. The full null replicate archive is reproducible from the saved scenario definitions and seeds but is not versioned because of its size. The manuscript figures are generated by `article/make_figures.py`; the decisive experiment is generated by `experiments/run_validation.py`; the post-hoc variance audit is `experiments/audit_variance_bias.py`. The baseline implementation is `src/welch_differential_mi/welch.py`.

Acknowledgements

This draft was prepared as part of a master’s research project. Formal supervisor, institutional, funding, conflict-of-interest, and author- contribution statements should be inserted before submission.

References

- [1] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3–4):379–423, 623–656, 1948. doi: 10.1002/j.1538-7305.1948.tb01338.x.
- [2] David R. Brillinger. Some data analyses using mutual information. *Brazilian Journal of Probability and Statistics*, 18(2):163–183, 2004.
- [3] Joseph T. Lizier. JIDT: An information-theoretic toolkit for studying the dynamics of complex systems. *Frontiers in Robotics and AI*, 1:11, 2014. doi: 10.3389/frobt.2014.00011.
- [4] Franklin E. Satterthwaite. An approximate distribution of estimates of variance components. *Biometrics Bulletin*, 2(6):110–114, 1946. doi: 10.2307/3002019.
- [5] Bernard L. Welch. The generalization of “student’s” problem when several different population variances are involved. *Biometrika*, 34(1–2):28–35, 1947. doi: 10.1093/biomet/34.1-2.28.
- [6] Kermit Hutcheson. A test for comparing diversities based on the shannon formula. *Journal of Theoretical Biology*, 29(1):151–154, 1970. doi: 10.1016/0022-5193(70)90124-4.
- [7] George A. Miller. Note on the bias of information estimates. In Henry Quastler, editor, *Information Theory in Psychology: Problems and Methods*, pages 95–100. Free Press, Glencoe, Illinois, 1955.
- [8] Roelof Moddemeijer. On estimation of entropy and mutual information of continuous distributions. *Signal Processing*, 16(3):233–248, 1989. doi: 10.1016/0165-1684(89)90132-1.
- [9] Liam Paninski. Estimation of entropy and mutual information. *Neural Computation*, 15(6):1191–1253, 2003. doi: 10.1162/089976603321780272.
- [10] Kirthivasan Kandasamy, Akshay Krishnamurthy, Barnabas Poczos, Larry Wasserman, and James M. Robins. Influence functions for machine learning: Nonparametric estimators for entropies, divergences and mutual informations. *arXiv preprint arXiv:1411.4342*, 2015. doi: 10.48550/arXiv.1411.4342.
- [11] Ricardo Mora and Javier Ruiz-Castillo. The statistical properties of the mutual information index of multigroup segregation. Working Paper 09-84, Universidad Carlos III de Madrid, Departamento de Economía, 2009.
- [12] Ricardo Mora and Javier Ruiz-Castillo. Entropy-based segregation indices. *Sociological Methodology*, 41(1):159–194, 2011. doi: 10.1111/j.1467-9531.2011.01237.x.
- [13] Rebecca Allen, Simon Burgess, Russell Davidson, and Frank Windmeijer. More reliable inference for the dissimilarity index of segregation. *The Econometrics Journal*, 18(1):40–66, 2015. doi: 10.1111/ectj.12039.
- [14] Corey B. Hart and Simon F. Giszter. A neural basis for motor primitives in the spinal cord. *Journal of Neuroscience*, 30(4):1322–1336, 2010. doi: 10.1523/JNEUROSCI.5894-08.2010.
- [15] Suman Sarkar and Biswajit Pandey. A study on the statistical significance of mutual information between morphology of a galaxy and its large-scale environment. *Monthly Notices of the Royal Astronomical Society*, 497(4):4077–4090, 2020. doi: 10.1093/mnras/staa2236.
- [16] Luke Y. Prince, Matthew M. Tran, Dorian Grey, Lydia Saad, Helen Chasiotis, Jeehyun Kwag,

- Michael M. Kohl, and Blake A. Richards. Neocortical inhibitory interneuron subtypes are differentially attuned to synchrony- and rate-coded information. *Communications Biology*, 4:935, 2021. doi: 10.1038/s42003-021-02437-y.
- [17] Andrew Steven Martin, Heather Guy, Michael Ray Gallagher, and III Neely, Ryan Reynolds. A guide to optimised spatiotemporal data co-location by mutual information maximisation. *Atmospheric Measurement Techniques*, 19:3511–3537, 2026. doi: 10.5194/amt-19-3511-2026.
- [18] Caroline M. Holmes and Ilya Nemenman. Estimation of mutual information for real-valued data with error bars and controlled bias. *Physical Review E*, 100(2):022404, 2019. doi: 10.1103/PhysRevE.100.022404.
- [19] Luke Mather, Elisabeth Oswald, Joe Bandenburg, and Marcin Wójcik. Does my device leak information? an a priori statistical power analysis of leakage detection tests. In *Advances in Cryptology – ASIACRYPT 2013*, pages 486–505. Springer, 2013. doi: 10.1007/978-3-642-42033-7_25.
- [20] Gabriel R. Palma, Silvio S. Zocchi, Wesley A. C. Godoy, and Jorge A. Wiendl. New confidence interval methods for shannon index. *arXiv preprint arXiv:2204.10073*, 2022. doi: 10.48550/arXiv.2204.10073.
- [21] EunYi Chung and Joseph P. Romano. Exact and asymptotically robust permutation tests. *Annals of Statistics*, 41(2):484–507, 2013. doi: 10.1214/13-AOS1090.
- [22] Alina A. von Davier. A corrected welch satterthwaite equation. *arXiv preprint arXiv:2602.20912*, 2026. doi: 10.48550/arXiv.2602.20912.